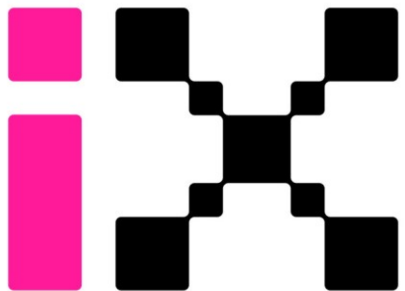


Differentiable simulators: a bridge between machine learning and scientific computing

Dr Aidan Crilly

Eric and Wendy Schmidt AI in Science postdoctoral fellow

Centre for Inertial Fusion Studies



Imperial-X

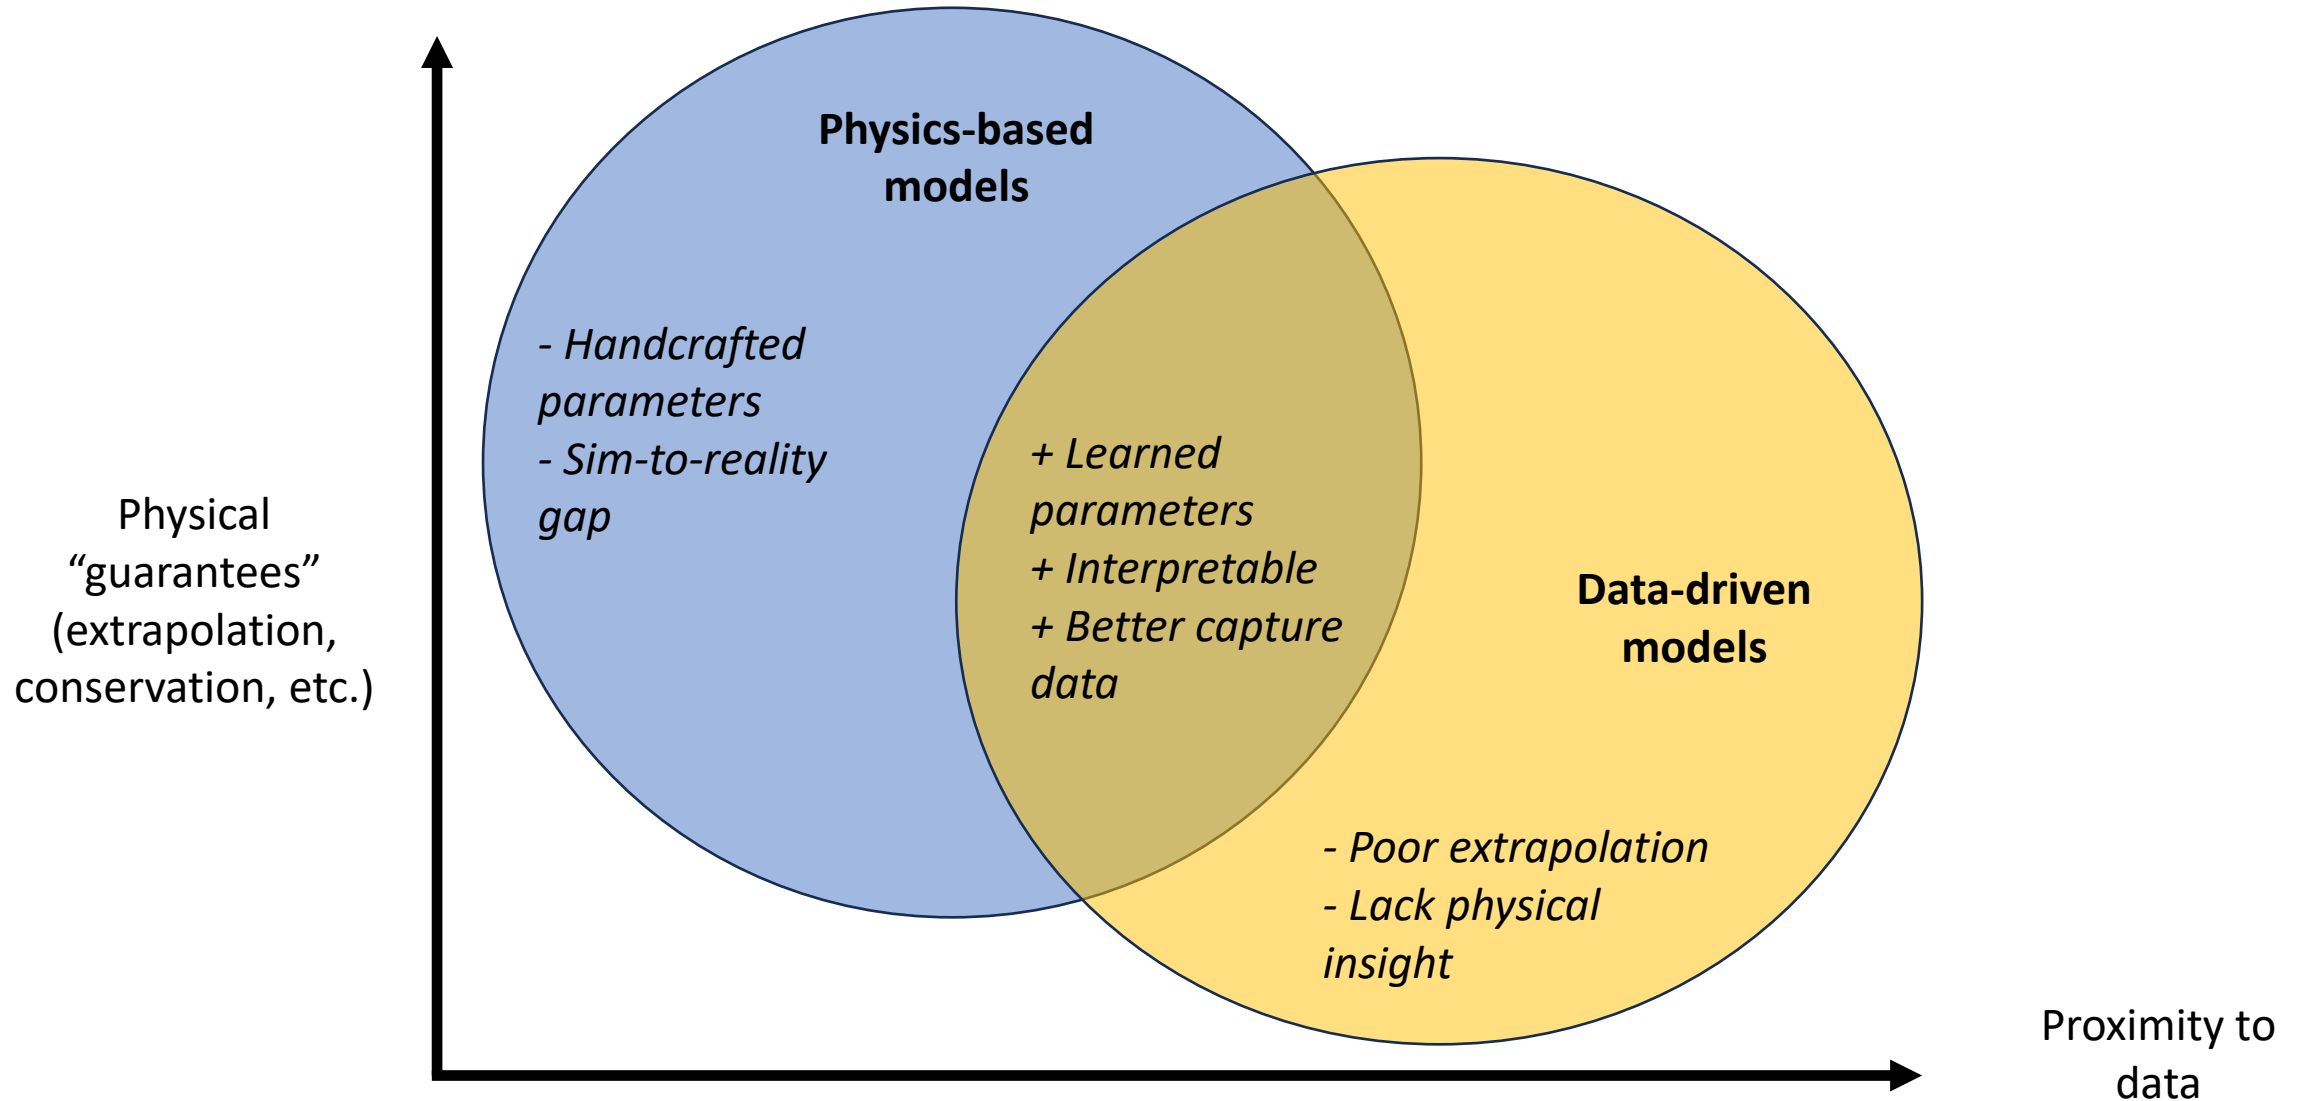


Outline

- Differential equations
 - Adjoint state problem
- Differentiable programming
 - For ODEs
 - For PDEs
- Gradient based optimisation
- Neural differential equations

*[https://github.com/
aidancrilly/MiniCourse-
DifferentiableSimulation](https://github.com/aidancrilly/MiniCourse-DifferentiableSimulation)*

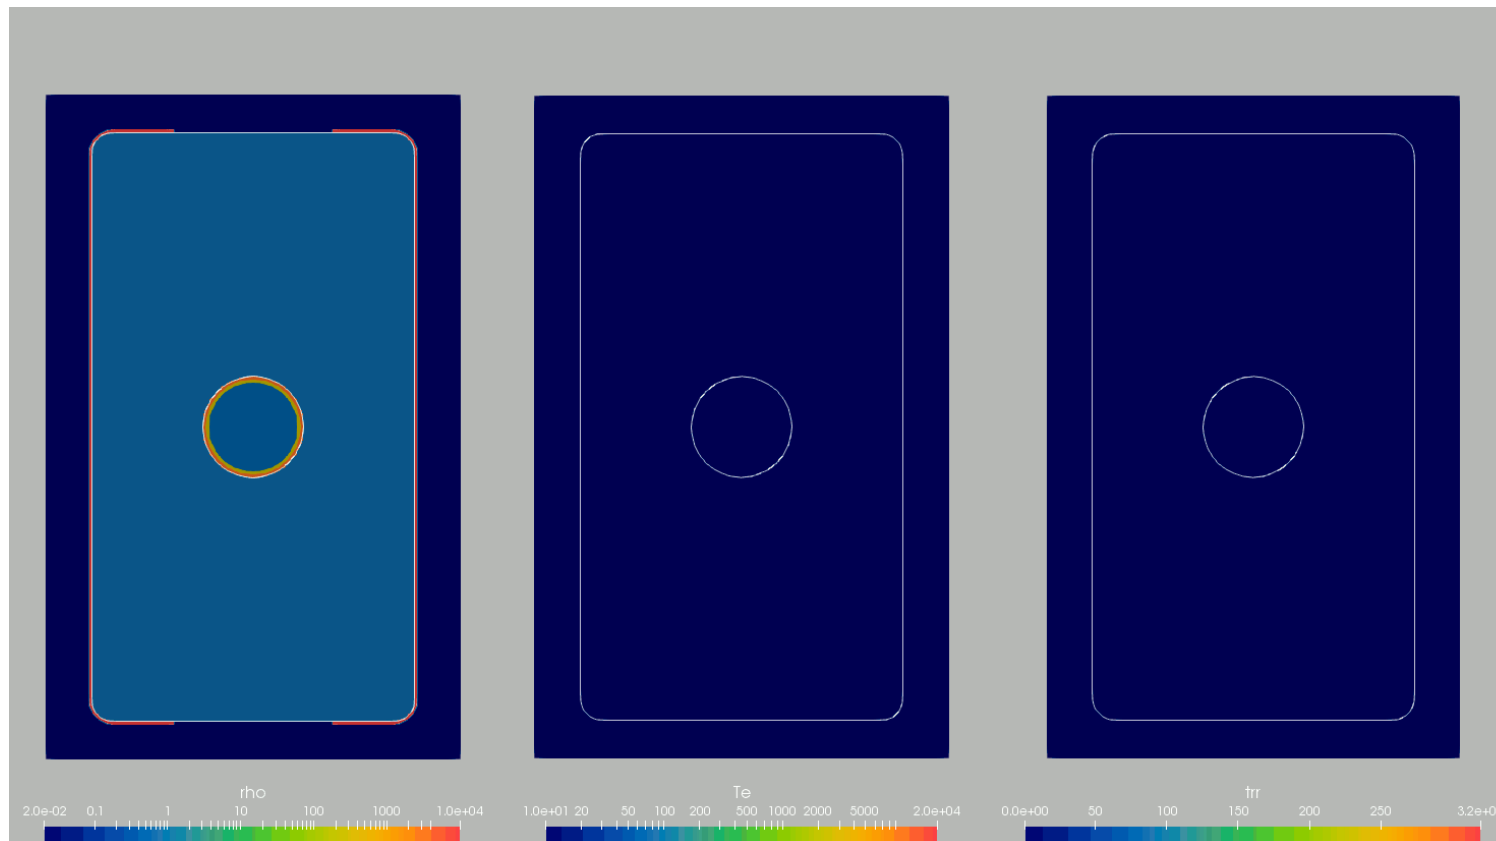
Data-driven and physics-based models



Context



- Multi-physics code in 3D:
 - Magneto-hydrodynamics, radiation transport, thermal conduction, split electron-ion energy equations, laser ray trace, alpha particle transport, non-ideal equation of state, extended Ohm's law, material strength...



Context



- Multi-physics code in 3D:
 - Magneto-hydrodynamics, radiation transport, thermal conduction, split electron-ion energy equations, laser ray trace, alpha particle transport, non-ideal equation of state, extended Ohm's law, material strength...

$$\partial_t \rho + \nabla \cdot (\rho \vec{u}) = 0$$

$$\partial_t (\rho \vec{u}) + \nabla \cdot (\rho \vec{u} \vec{u}) = -\nabla P + J \times B$$

$$\partial_t (\rho e_{tot}) + \nabla \cdot (\rho e_{tot} \vec{u}) = -\nabla \cdot (P \vec{u}) + \text{Sources}$$

$$\frac{\partial \underline{B}}{\partial t} = \nabla \times \underline{v}_B \times \underline{B} - \nabla \times \eta_{\parallel} \nabla \times \underline{B} + \nabla \times \frac{\nabla P_e}{en_e}$$

$$\left[\frac{\partial}{\partial t} + \kappa_{\nu} c \right] E_{\nu} = 4\pi j_{\nu} - \nabla \cdot \vec{F}_{\nu} ,$$

$$\left[\frac{\partial}{\partial t} + \kappa_{\nu} c \right] \vec{F}_{\nu} = -c^2 \nabla \cdot \mathbf{P}_{\nu} ,$$

$$C_e \frac{\partial T_e}{\partial t} = \kappa_e \nabla^2 T_e + S_e + \omega_{ie} (T_i - T_e)$$

$$C_i \frac{\partial T_i}{\partial t} = \kappa_i \nabla^2 T_i + S_i + \omega_{ie} (T_e - T_i)$$

$$\frac{\partial v_{ray}}{\partial t} = \nabla \cdot \left(\frac{-c^2}{2} \frac{n_e}{n_c} \right)$$

$$\begin{aligned} & \left[\frac{1}{v} \frac{\partial}{\partial t} + \hat{\Omega} \cdot \vec{\nabla} + n(\vec{r}, t) \sigma(E) \right] \psi(\vec{r}, \hat{\Omega}, E, t) \\ & = S_{ex}(\vec{r}, \hat{\Omega}, E, t) \\ & + \int_0^{\infty} dE' \int d\hat{\Omega}' n(\vec{r}, t) \sigma_s(\hat{\Omega}' \cdot \hat{\Omega}, E' \rightarrow E) \psi(\vec{r}, \hat{\Omega}', E', t), \end{aligned}$$

$$\begin{aligned} \frac{\partial f_{\alpha}(v_{\alpha}, t)}{\partial t} &= C(f_{\alpha}(v_{\alpha}, t)) + \frac{S_0 \delta(v_{\alpha} - v_b)}{4\pi v_{\alpha}^2} \\ &= \frac{1}{\tau_S v_{\alpha}^2} \frac{\partial}{\partial v_{\alpha}} \left\{ \frac{v_{\alpha}^3 T_e + v_c^3 T_i \Phi(x_i)}{m_{\alpha} v_{\alpha}} \frac{\partial f_{\alpha}}{\partial v_{\alpha}} \right. \\ & \quad \left. + [v_{\alpha}^3 + v_c^3 \Phi(x_i)] f_{\alpha} \right\} + \frac{S_0 \delta(v_{\alpha} - v_b)}{4\pi v_{\alpha}^2}, \quad \text{And the list goes on...} \end{aligned}$$

The transport equation

$$\frac{\partial f}{\partial t} + \boxed{\mathbf{v} \cdot \nabla f} + \frac{1}{m} \boxed{\mathbf{F} \cdot \nabla_{\mathbf{v}} f} = \boxed{\left(\frac{\partial f}{\partial t} \right)_{coll.}} + \boxed{\left(\frac{\partial f}{\partial t} \right)_{src}}$$

Advection **Forcing** **Collisions** **Sources**

- Encodes conservation laws on particle number, momenta and energy
- 7-dimensional (3-space, 3-velocity, 1-time), $f(x, y, z, v_x, v_y, v_z, t)$

The closure problem

- N^{th} moment equation depends on the $N+1^{\text{th}}$ moment (or more)

For example:

Integrating out velocity space:

$$\frac{\partial n}{\partial t} + \frac{\partial}{\partial x} (n \mathbf{V}) = 0$$

Multiply by velocity, and integrate out velocity space:

$$\frac{\partial n \mathbf{V}}{\partial t} + \frac{\partial}{\partial x} \left(n \mathbf{V} \mathbf{V} + \frac{1}{m} \mathbf{P} \right) = 0$$

Traditional approaches

- Moment 'ordering' argument
 - "The $N+1$ 'th (and higher orders) are negligibly small"
- Proximity to equilibrium
 - "The steady state distribution can be used to compute the closure"
- Limiting behaviour
 - "I expect this limiting behaviour, I invent a closure which respects this"
- Asymptotic expansion
 - "I will expand in terms of a small parameter and neglect higher order terms"

The King of Transport

- Flux limited diffusion is ubiquitous
- In general, this looks like:

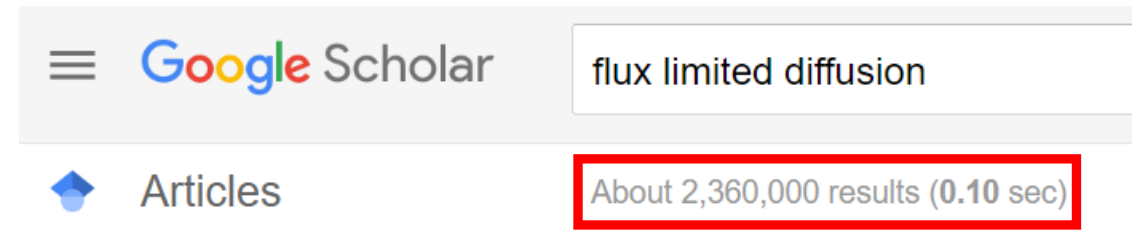
Diffusion equation: $\frac{\partial E}{\partial t} + \nabla \cdot F = Q$

Diffusivity

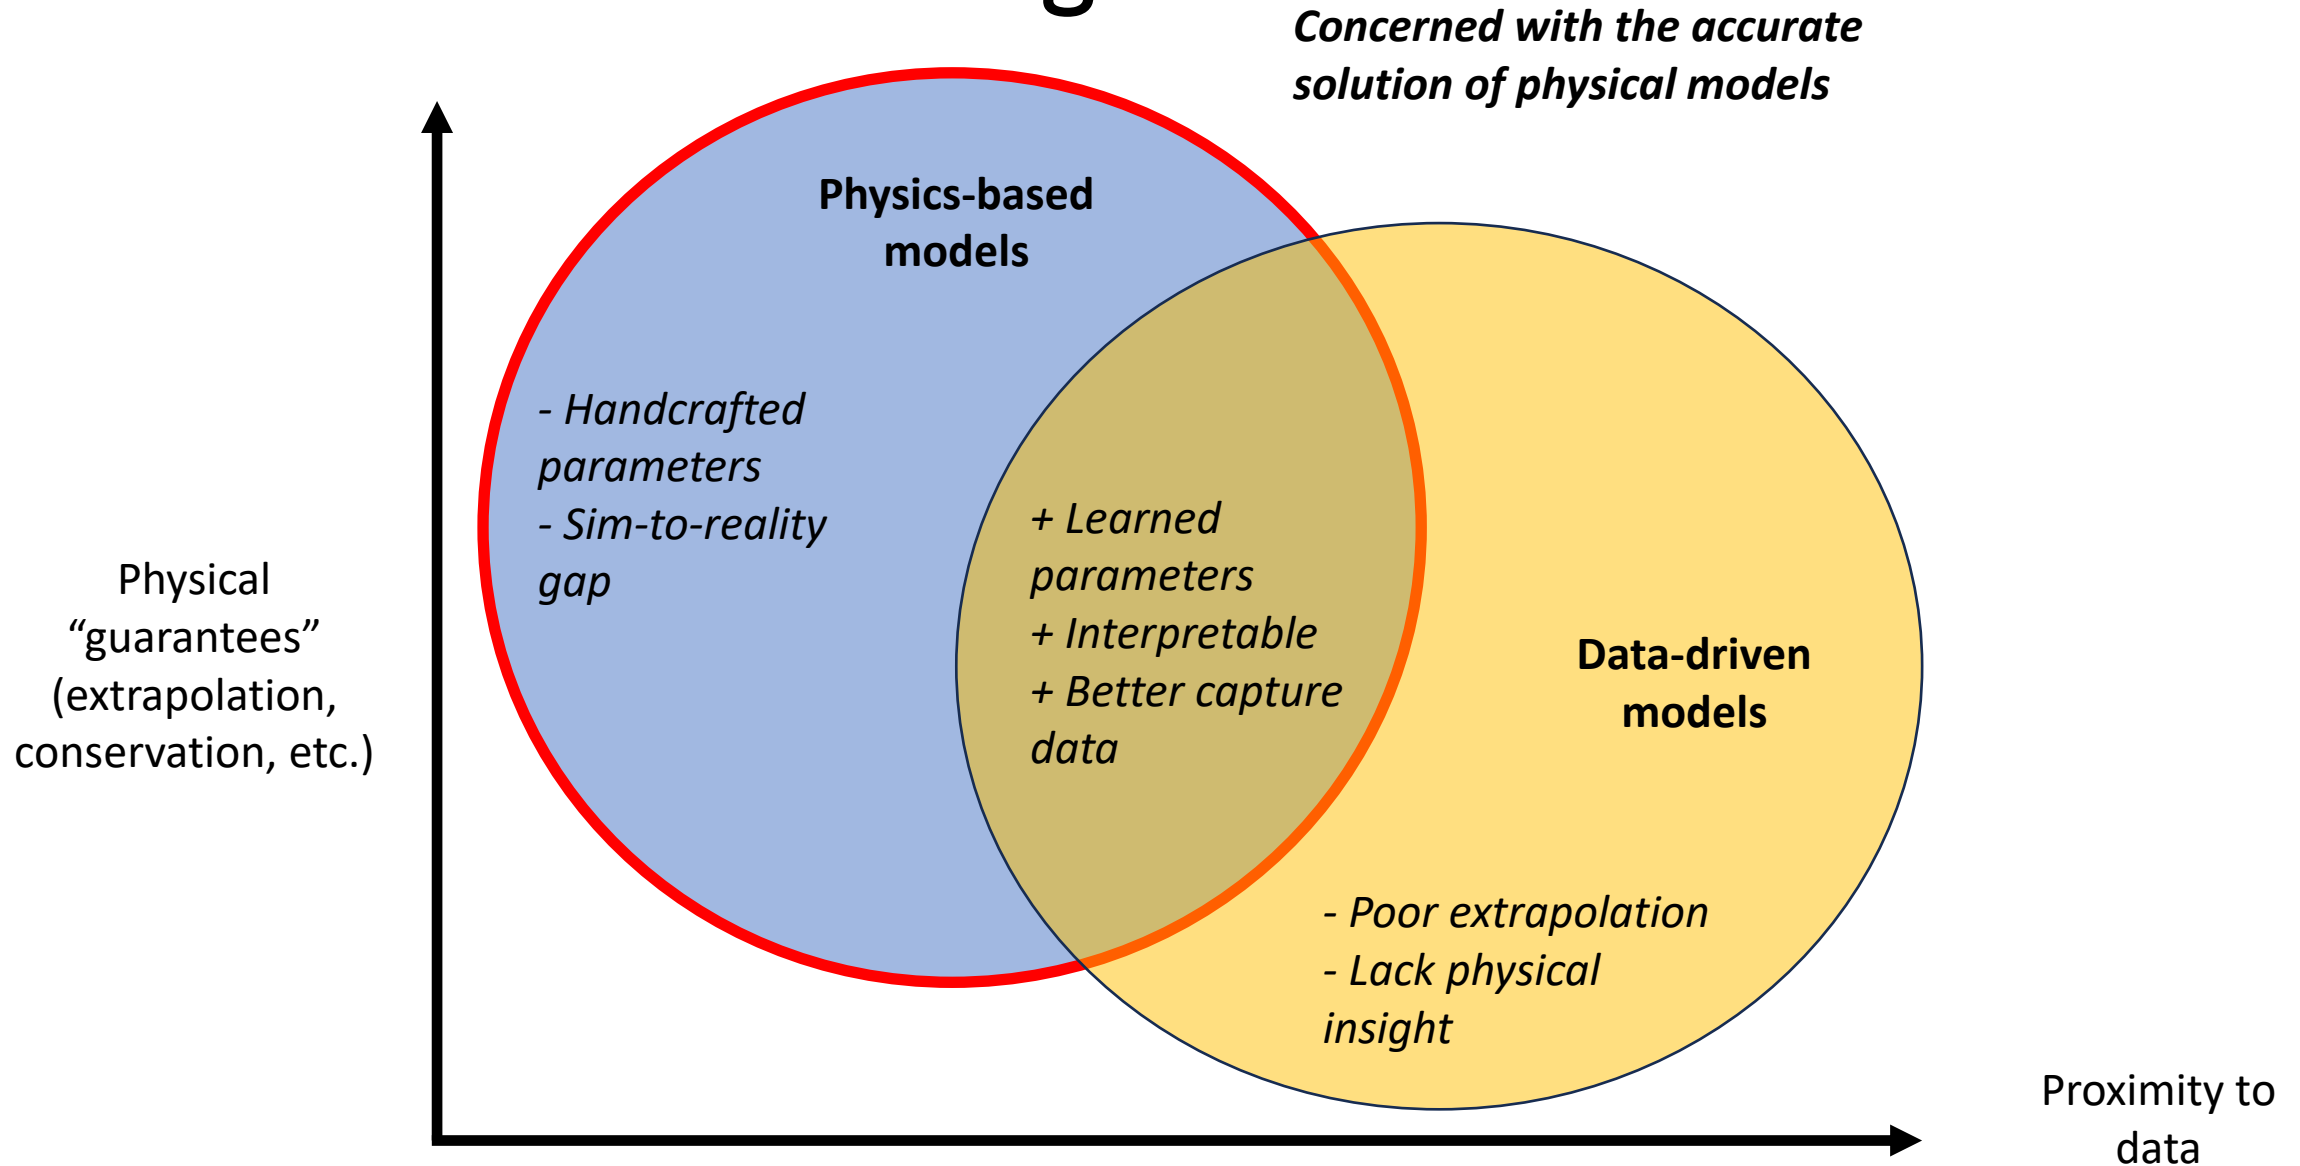
Flux closure: $F = -\lambda D \nabla E$

Flux limiter

- Flux limiters can be *incredibly* ad-hoc – ruins the predictive power of simulation codes



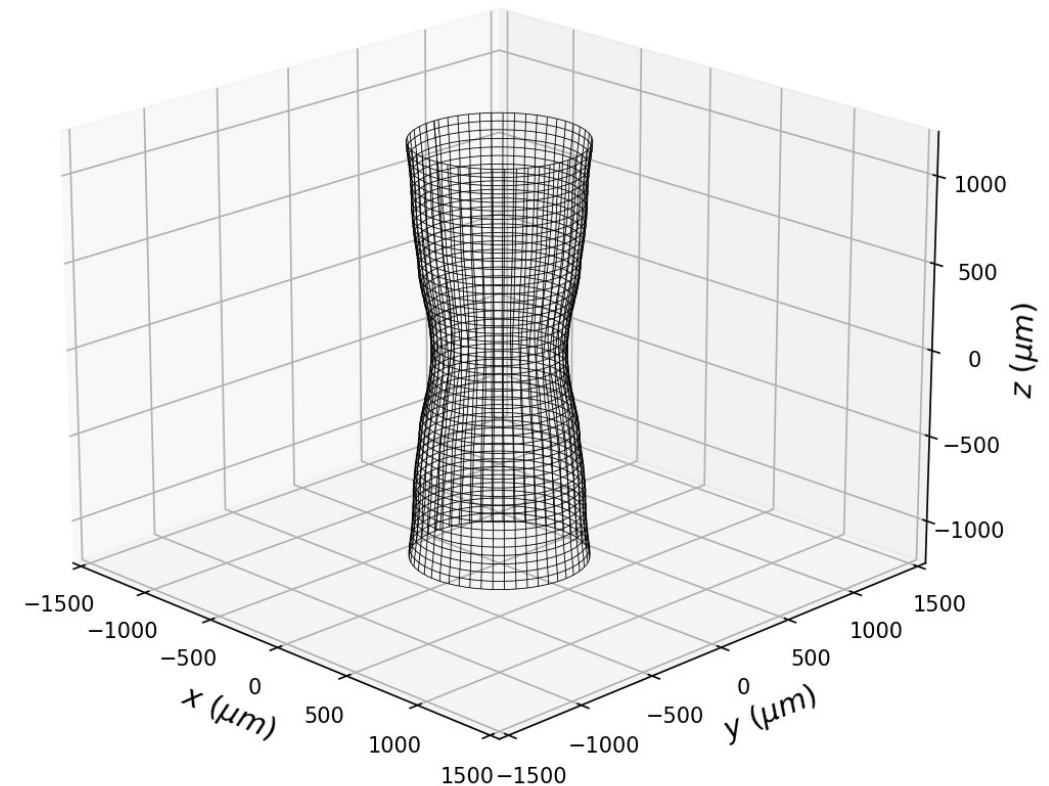
Numerical modelling



Numerical modelling – ODEs

- Ordinary differential equations (ODEs) – dependent on only a single independent variable
 - For the most part, this variable is time
- Example: the path of a laser in a plasma (geometric optics)

$$\frac{d\underline{v}_{ray}}{dt} = \nabla \left(\frac{-c^2}{2} \frac{n_e}{n_c} \right)$$



Numerical modelling – ODEs

- Numerical solutions to ODEs use finite steps (h or dt) to approximate derivatives

$$\frac{dy}{dt} = f(t, y)$$

Numerical modelling – ODEs

- We use finite differencing to approximate derivatives

Forward differencing:

$$\frac{dy}{dx} \approx \frac{\Delta y}{\Delta x} = \frac{y^{n+1} - y^n}{x^{n+1} - x^n}$$

Backward differencing:

$$\frac{dy}{dx} \approx \frac{\Delta y}{\Delta x} = \frac{y^n - y^{n-1}}{x^n - x^{n-1}}$$

First order in accuracy

Centred differencing:

$$\frac{dy}{dx} \approx \frac{\Delta y}{\Delta x} = \frac{y^{n+1} - y^{n-1}}{x^{n+1} - x^{n-1}} \longrightarrow$$

Second order in accuracy

Numerical modelling – ODEs

- Numerical solutions to ODEs use finite steps (h or dt) to approximate derivatives

$$\frac{dy}{dt} = f(t, y)$$

- Simplest = Forward Euler:

$$y_{n+1} = y_n + h f(t_n, y_n)$$

Numerical modelling – ODEs

- Numerical solutions to ODEs use finite steps (h or dt) to approximate derivatives

$$\frac{dy}{dt} = f(t, y)$$

- Simplest = Forward Euler:

$$y_{n+1} = y_n + h f(t_n, y_n)$$

- Very common = 4th order Runge-Kutta (RK4) with adaptive stepping

$$y_{n+1} = y_n + \frac{h}{6} (k_1 + 2k_2 + 2k_3 + k_4),$$

$$t_{n+1} = t_n + h$$

$$k_1 = f(t_n, y_n),$$

$$k_2 = f\left(t_n + \frac{h}{2}, y_n + h \frac{k_1}{2}\right),$$

$$k_3 = f\left(t_n + \frac{h}{2}, y_n + h \frac{k_2}{2}\right),$$

$$k_4 = f(t_n + h, y_n + h k_3).$$

Forward vs Inverse problems

- Solutions to differential equations are often concerned with the *forward problem*

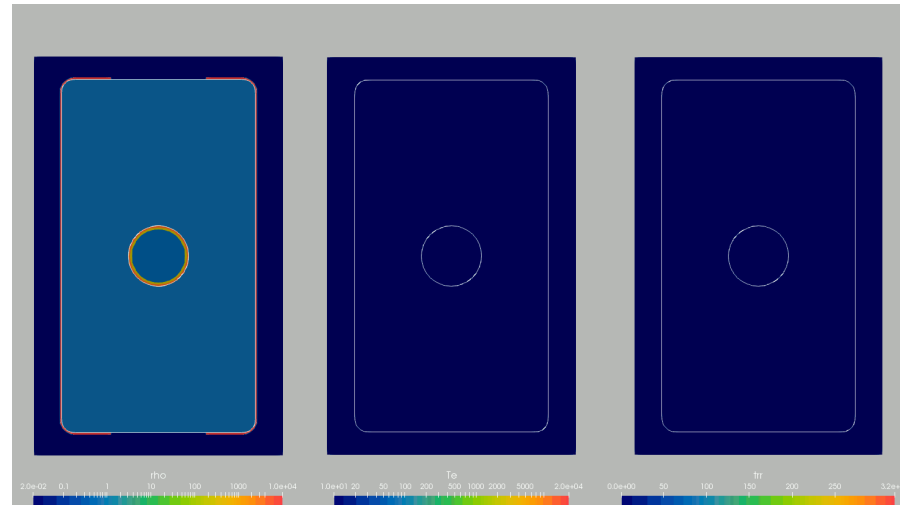
Differential equation: $h(x, p, t) = 0$
Initial condition: $g(x(0), p) = 0$

x = state variables
 p = parameters
 t = time

Initial conditions:
Densities, temperatures, etc.

Differential equations:
Hydrodynamics ++

Other parameters:
Laser power vs time



Forward vs Inverse problems

- Solutions to differential equations are often concerned with the *forward problem*

Differential equation: $h(x, p, t) = 0$

Initial condition: $g(x(0), p) = 0$

x = state variables

p = parameters

t = time

- However, if we want to minimize some other scalar function at the same time = *inverse problem*

Figure of merit: Minimise w. r. t. p : $f(x, p)$

Forward vs Inverse problems

- Solutions to differential equations are often concerned with the *forward problem*

Differential equation: $h(x, p, t) = 0$

Initial condition: $g(x(0), p) = 0$

x = state variables

p = parameters

t = time

- However, if we want to minimize some other scalar function at the same time = *inverse problem*

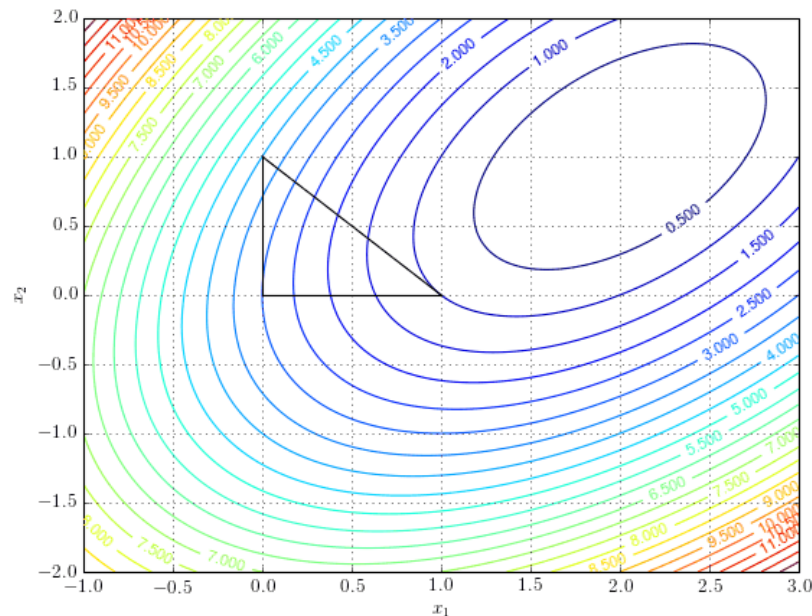
Figure of merit: **Minimise w. r. t. p : $f(x, p)$**

Optimisation

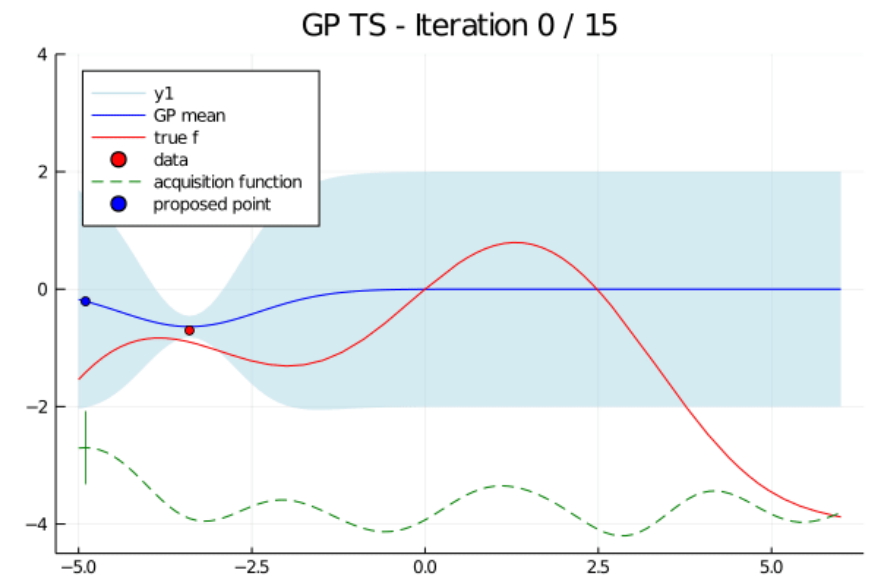
Optimisation

- Minimisation = optimisation = root-finding of gradient
- Gradient-based vs gradient-free optimisation
- Gradient-free:

Downhill-simplex

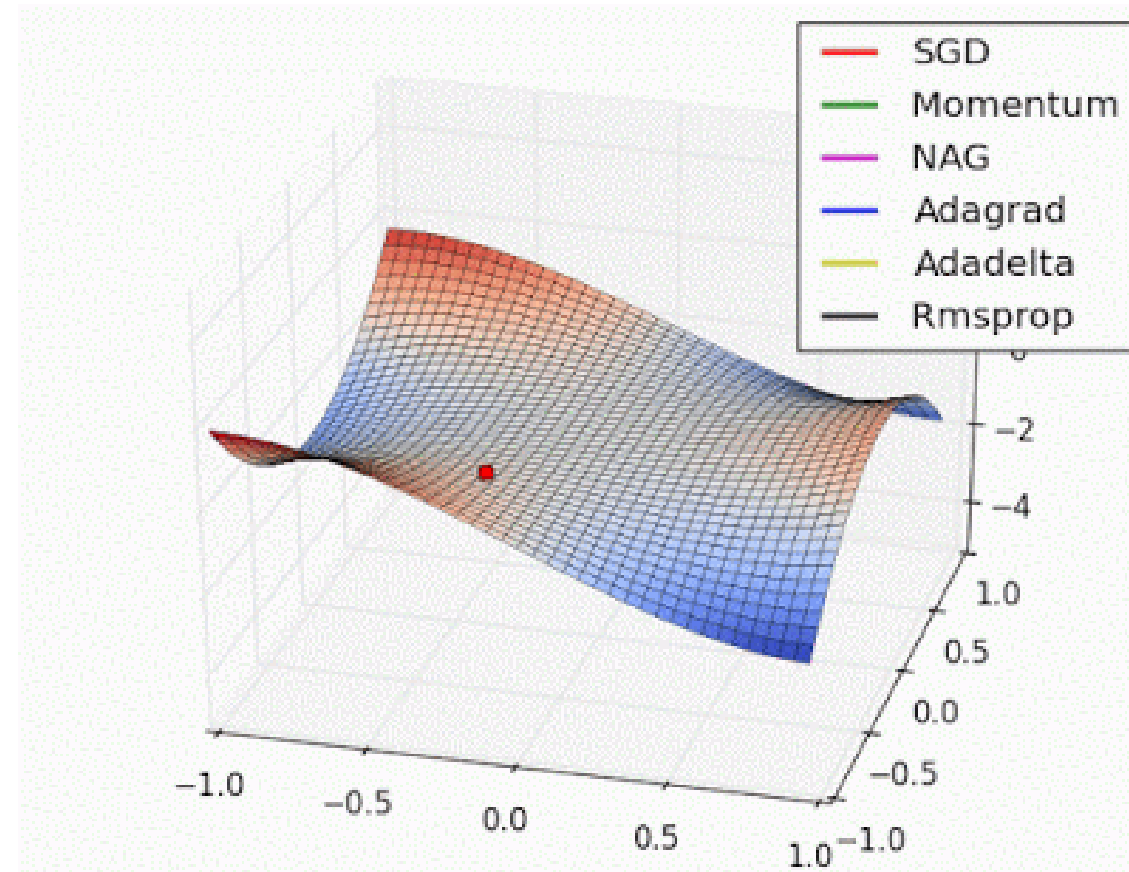


Bayesian optimisation



Optimisation

- Minimisation = optimisation
- Gradient-based:
 - Simple idea, roll down-hill (to minimize)
- We will visit the mathematics later...

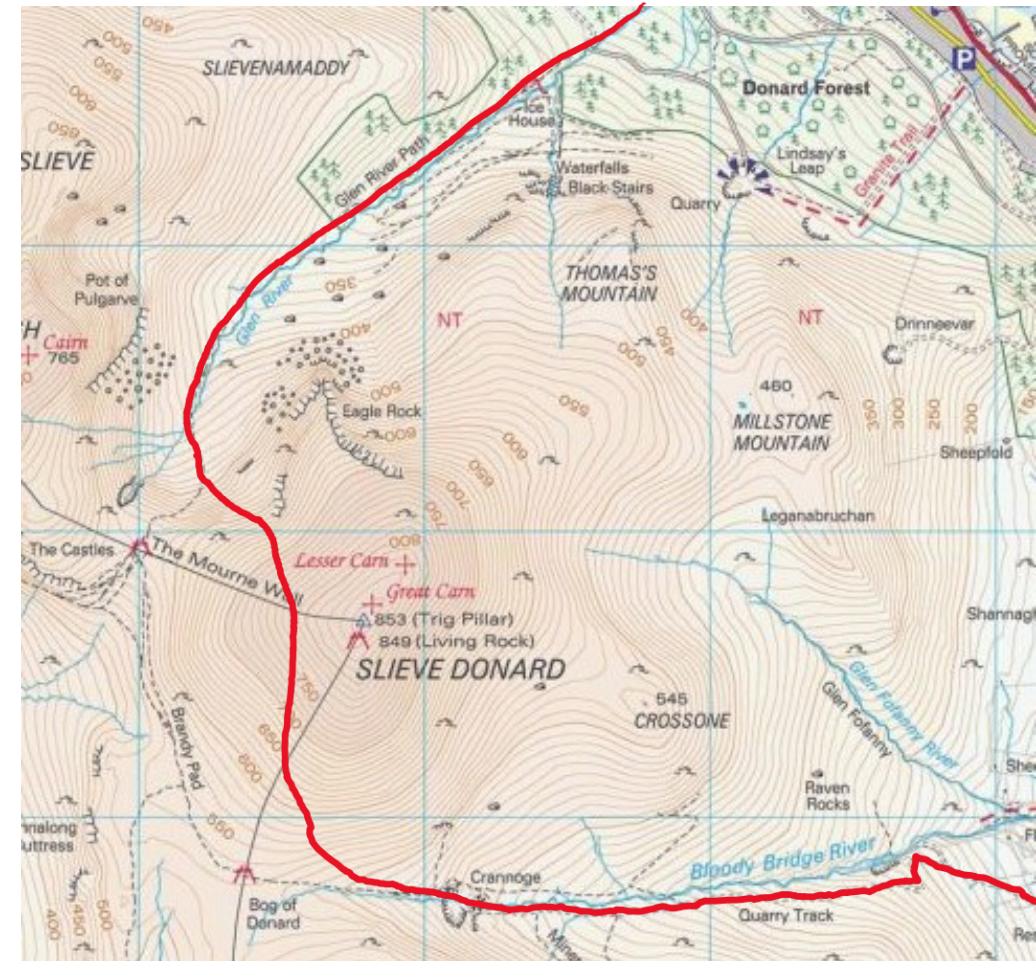


The adjoint state problem

- How do we solve the inverse problem in differential equations efficiently?
- We could run additional forward solutions to get finite difference gradients – *very inefficient*
- Instead of predicting how a *single* design change influences *every* aspect of the flow, the adjoint method predicts how *every* design change influences a *single* aspect of the flow.

Constrained Optimisation

- Inverse problem in simplified form:
Maximise w.r.t. p : $f(p)$
 $g(p) = 0$
- *Intuitive example*: Reach highest point (f) while sticking to path (g)
 - When path falls below you both in front and behind you
 - Or the path is tangent with contour
 - Or the path of steepest ascent (gradient) is 90 degree to path



In this example:

- $f(p)$ gives the height given your position, p , i.e. x and y coordinates.
- $g(p)$ parameterises the path you take, such that if you are on the path $g(p) = 0$

Constrained Optimisation – the maths

- Combine target and constraint using “Lagrange multipliers”



$$L(p, \lambda) = f(p) + \lambda g(p)$$

- Find optima of this function:

$$\partial_p L = 0 \rightarrow \partial_p f = -\lambda \partial_p g$$

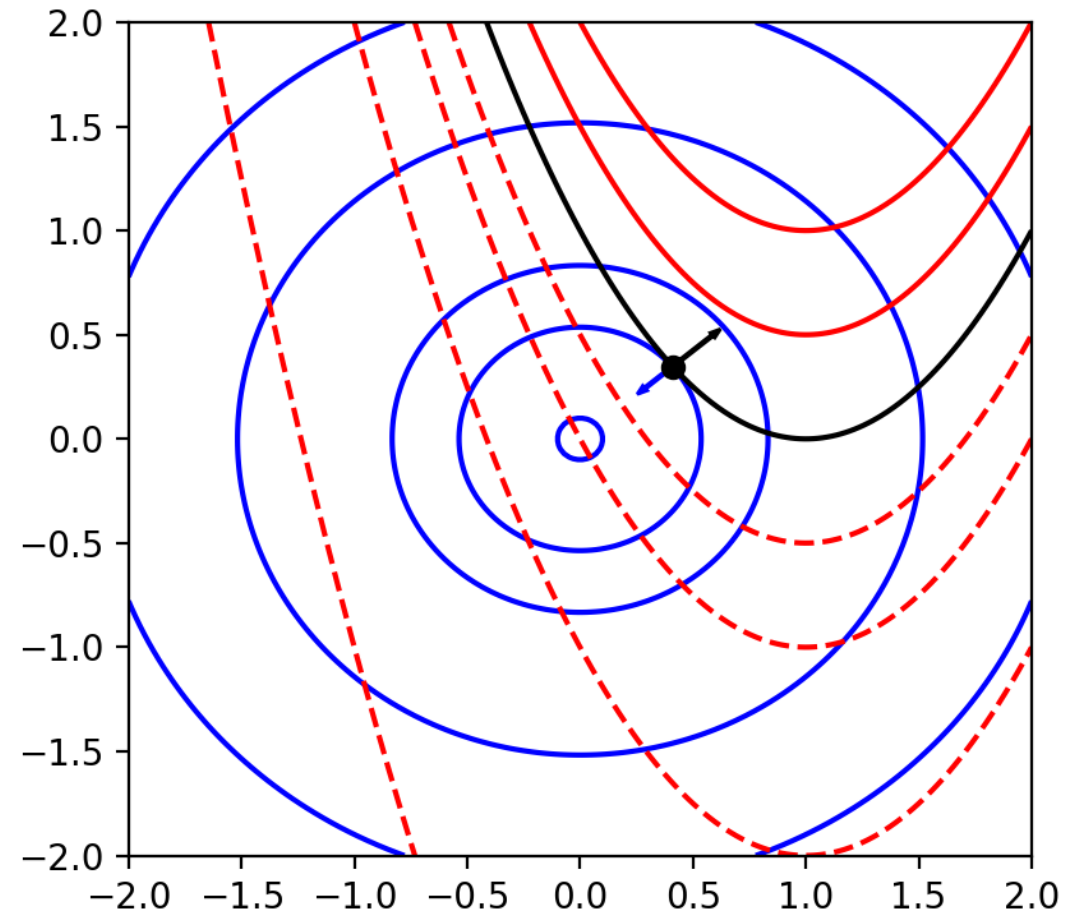
$$\partial_\lambda L = g(p) = 0$$

- In other words, at the optima the gradients are parallel

Let's consider $p = [x, y]$ and:

$$f(x, y) = e^{-x^2 - y^2}$$

$$g(x, y) = y - (x - 1)^2 = 0$$



The adjoint state problem

- Inverse problem in simplified form:

$$\begin{aligned} &\text{Minimise w.r.t. } p: f(x, p) \\ &g(x, p) = 0 \end{aligned}$$

- Constrained optimisation \rightarrow Lagrange multipliers:

$$L(x, p, \lambda) = f(x, p) + \lambda^T g(x, p)$$

- Gives *adjoint equation* and gradient w.r.t. parameters, p :

$$\frac{\partial f}{\partial x} + \lambda^T \frac{\partial g}{\partial x} = 0, \quad \frac{df}{dp} = \frac{dL}{dp} = \frac{\partial f}{\partial p} + \lambda^T \frac{\partial g}{\partial p}$$

The adjoint state problem – for ODEs

- A simple example ODE, with parameters y_0 and τ

$$\frac{dy}{dt} = -\frac{1}{\tau}y, \quad f = L(y(T))$$

- What are the formal *adjoint equations*?
 - There will be differential equations for
 1. *Adjoint process (Lagrange multiplier)*
 2. *Parameter gradients (df/dp)*

The adjoint state problem – for ODEs

- To the whiteboard!

What if our function is a solution to an ODE?

- A simple example ODE, with parameters y_0 and τ

$$\frac{dy}{dt} = -\frac{1}{\tau}y, \quad f = L(y(T))$$

- What are the formal *adjoint equations*?

Adjoint process: $\frac{d\alpha}{dt} = \frac{1}{\tau}\alpha, \quad \alpha(T) = \frac{dL}{dy(T)}$ Exponential growth

Parameter gradients: $\frac{d\beta}{dt} = -\frac{1}{\tau^2}\alpha y, \quad \beta(T) = 0$

- Solved backwards in time, $\beta(0) = d_{\tau}f = \left(\frac{T}{\tau^2}\right) \left(\frac{dL}{dy(T)}\right) y_0 \exp\left[-\frac{T}{\tau}\right]$

Automatic differentiation introduction

- Numerical differentiation, finite difference methods:

$$\frac{dy}{dt} = \lim_{\varepsilon \rightarrow 0} \frac{y(t + \varepsilon) - y(t)}{\varepsilon} = f(t, y)$$

- Numerical error dependent on the step size used:

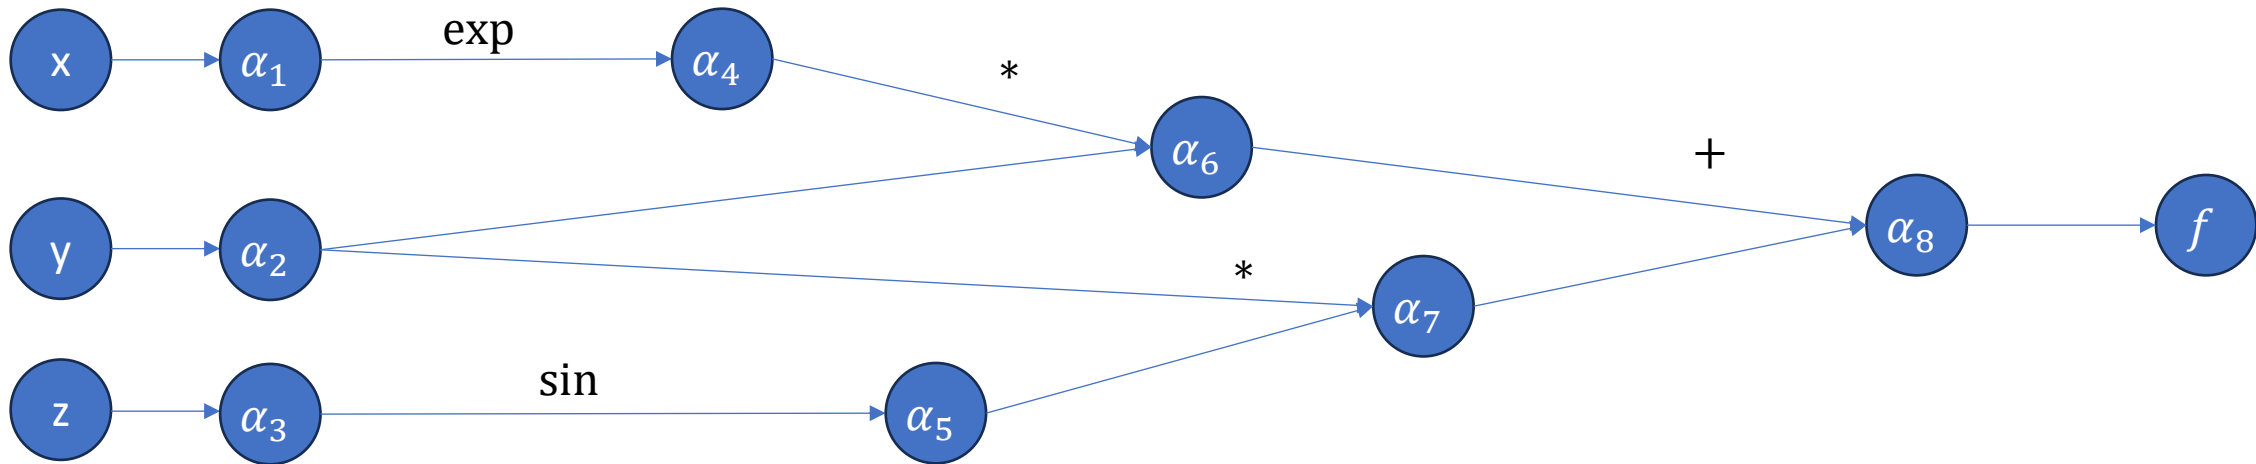
$$\frac{y(t + \varepsilon) - y(t)}{\varepsilon} = \dot{y}(t) + \frac{1}{2} \varepsilon \ddot{y}(t) + \dots = f(t, y) + \text{Error}(\varepsilon)$$

- Enter automatic differentiation...

Automatic Differentiation Intro

- Every computer program can be written as a computation graph of mathematical 'primitive' functions ($+$, \times , \exp , \sin , etc.)
- For example:

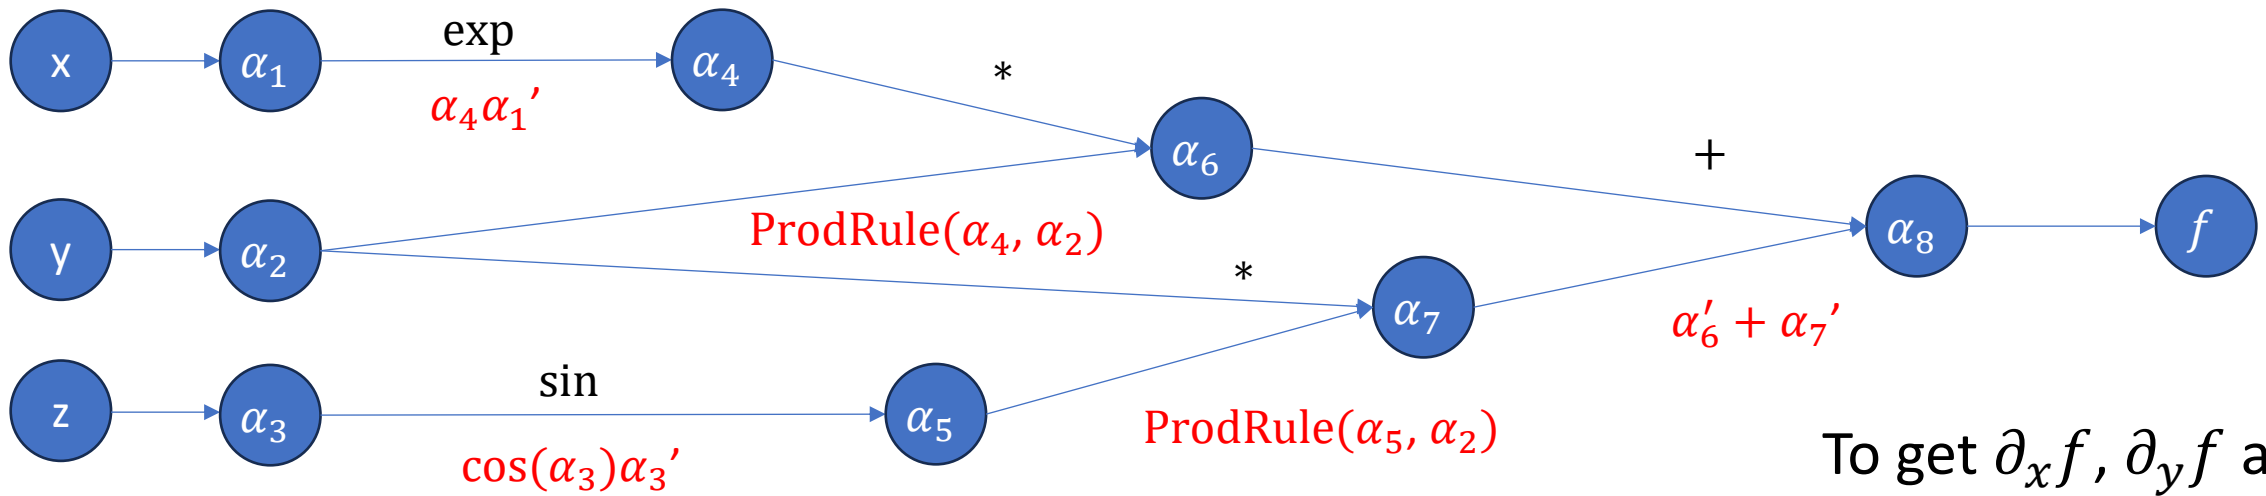
$$f(x,y,z) = y \exp(x) + y \sin(z)$$




Automatic Differentiation Intro

- Every computer program can be written as a computation graph of mathematical ‘primitive’ functions ($+$, \times , \exp , \sin , etc.)
- For example:

$$f(x,y,z) = y \exp(x) + y \sin(z)$$




“Forward mode” automatic differentiation

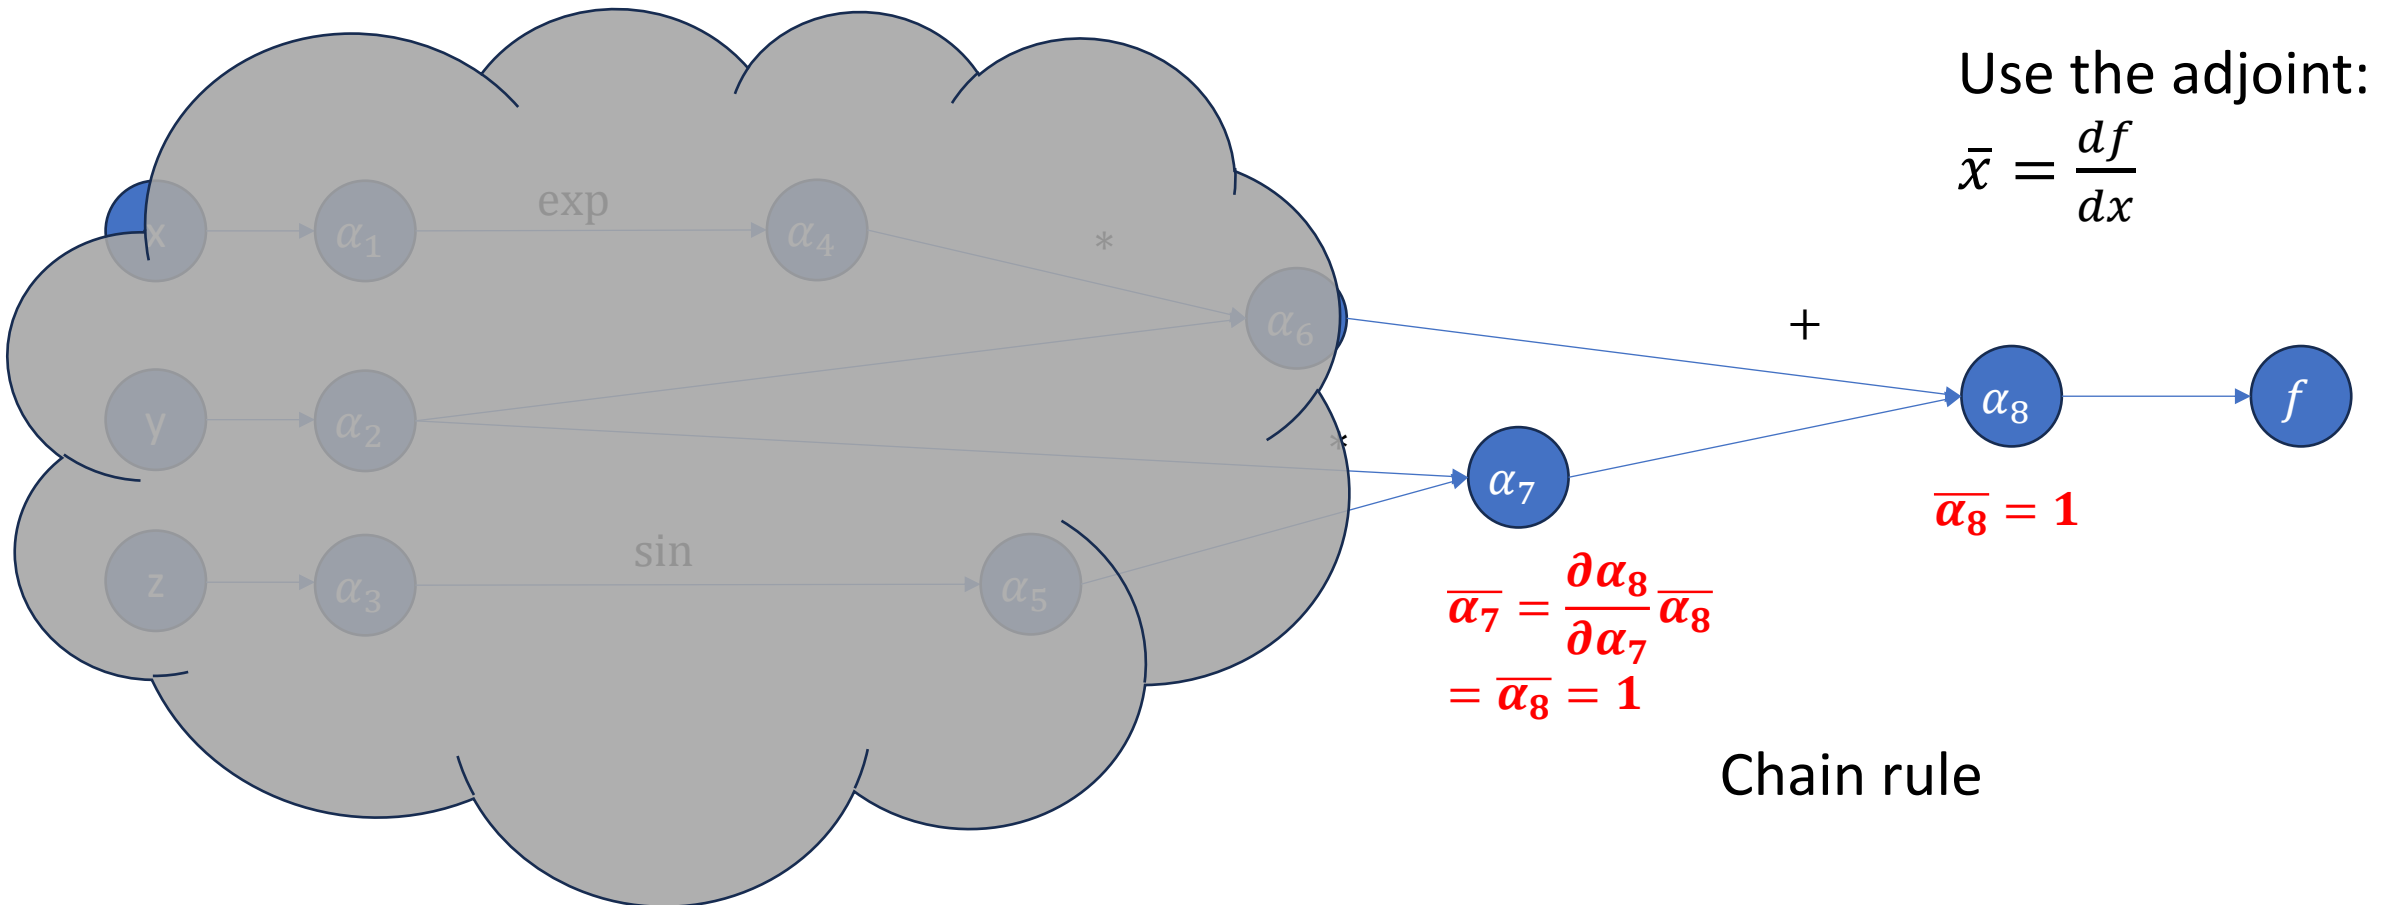
To get $\partial_x f$, $\partial_y f$ and $\partial_z f$
takes 3 passes through
graph

Automatic Differentiation Intro

- All optimisation problems have many-to-one functions
- Reverse-mode AD provides most efficient method to compute gradients
- Define the 'adjoint' as gradient w.r.t. output: $\bar{x} = \frac{df}{dx}$

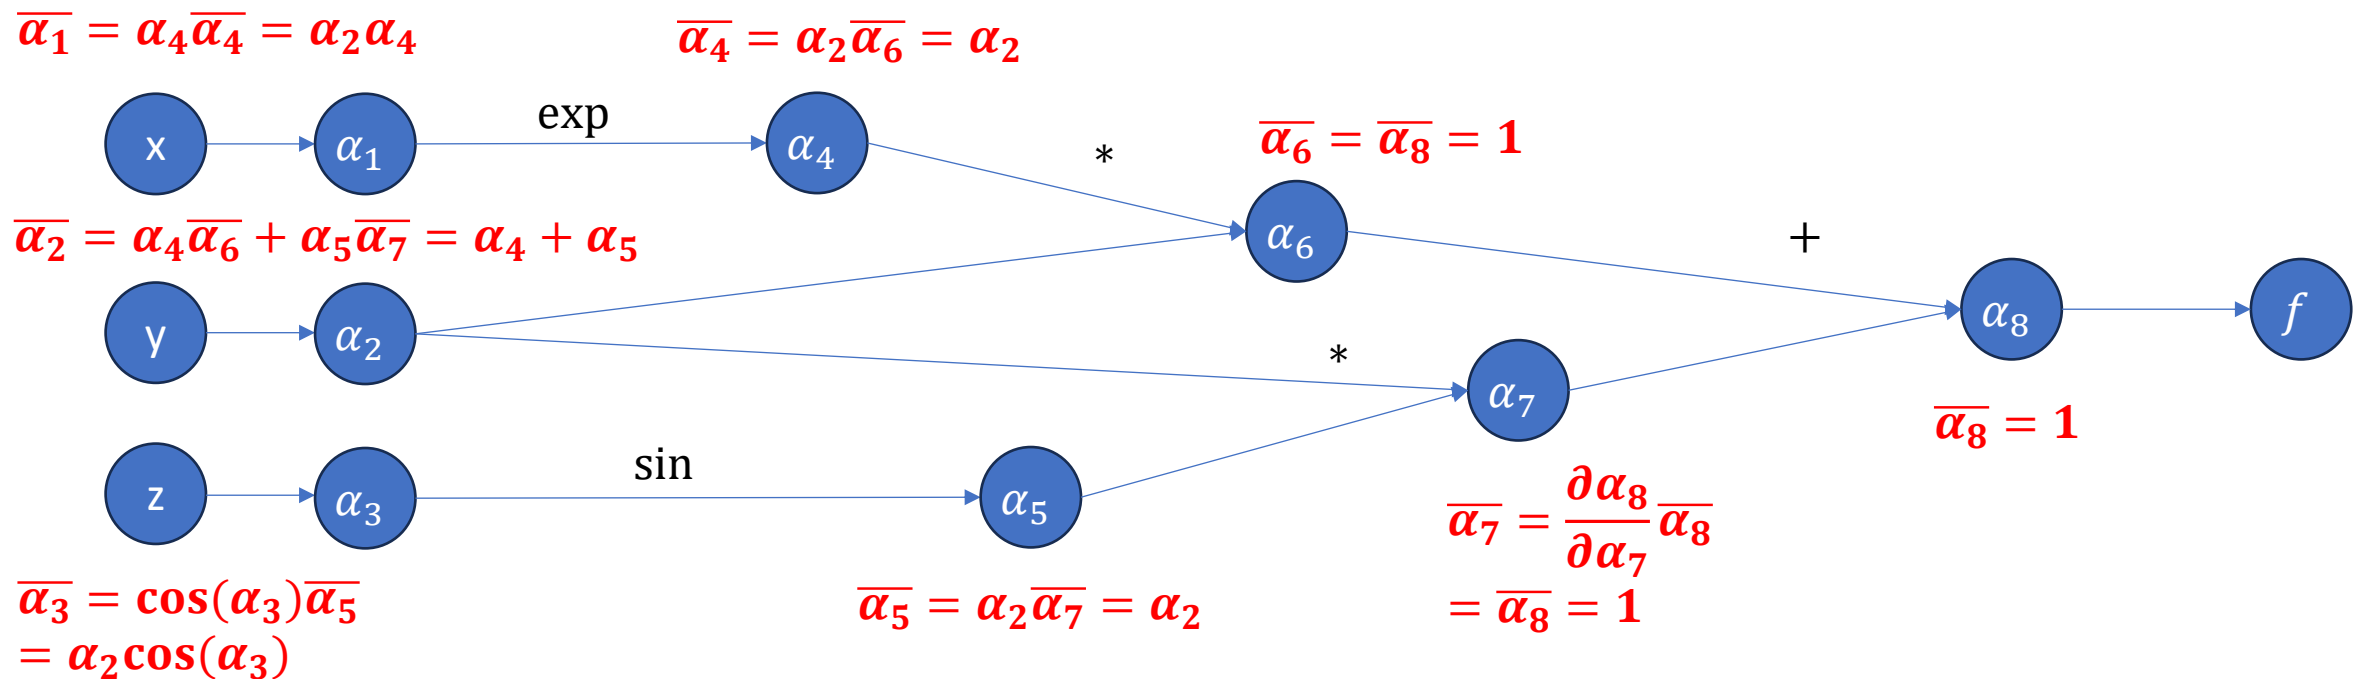
Automatic Differentiation Intro

- In our example:



Automatic Differentiation Intro

- In our example:



A single 'back-propagation' gives *all* gradients!

Differentiable Programming?



Yann LeCun:

“Deep Learning est mort. Vive Differentiable Programming!

...

*An increasingly large number of people are defining the networks procedurally in a data-dependent way (with loops and conditionals), allowing them to change dynamically as a function of the input data fed to them. **It's really very much like a regular program, except it's parameterized, automatically differentiated, and trainable/optimizable.** Dynamic networks have become increasingly popular...”*

Automatic-differentiation/differentiable-programming frameworks

- Python: PyTorch – autograd
- Python: Tensorflow – GradientTape
- Python: JAX
- C++: autodiff
- Check out autodiff.org for other languages/libraries

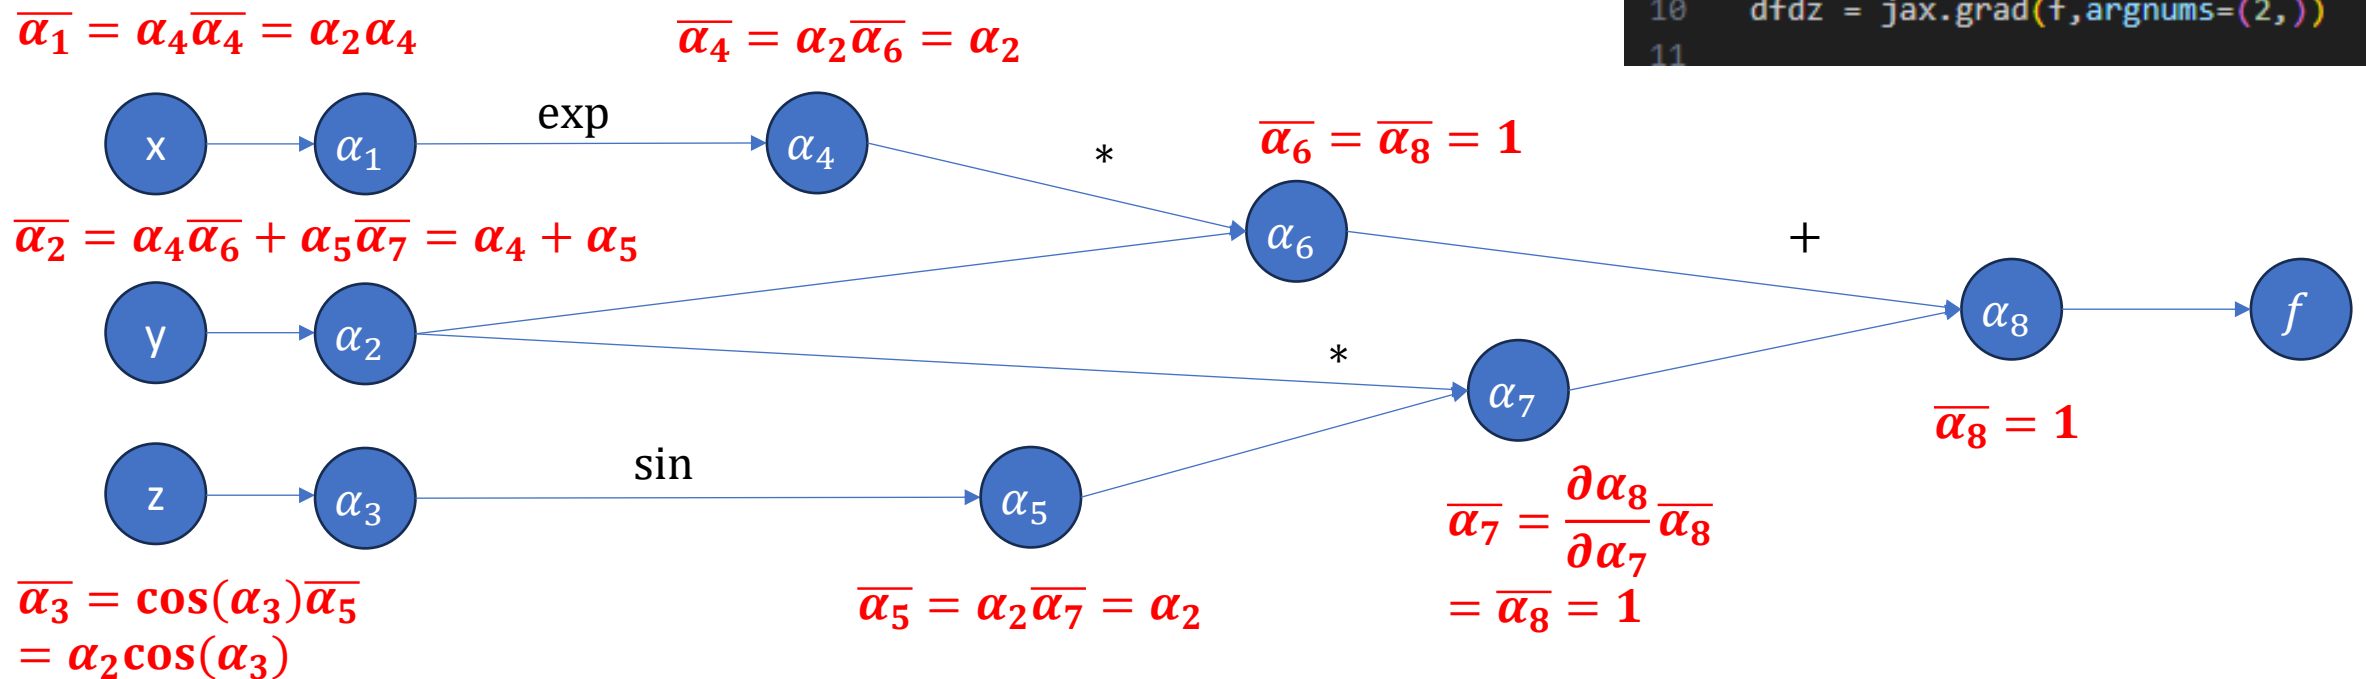
Automatic-differentiation/differentiable-programming frameworks

- Python: PyTorch – autograd
- Python: Tensorflow – GradientTape
- **Python: JAX**
- C++: autodiff
- Check out autodiff.org for other languages/libraries

PythonJAX simple example

- In example:
 - $f(x, y, z) = y \exp(x) - y \sin(z)$

```
1 import jax
2 import jax.numpy as jnp
3 import matplotlib.pyplot as plt
4
5 def f(x,y,z):
6     return y*jnp.exp(x)-y*jnp.sin(z)
7
8 dfdx = jax.grad(f,argnums=(0,))
9 dfdy = jax.grad(f,argnums=(1,))
10 dfdz = jax.grad(f,argnums=(2,))
11
```



A single 'back-propagation' gives *all* gradients!

Colab exercise 0

- An introduction to JAX

*[https://github.com/
aidancrilly/MiniCourse-
DifferentiableSimulation](https://github.com/aidancrilly/MiniCourse-DifferentiableSimulation)*

Back to differential equations...

The adjoint state problem – differentiable programming

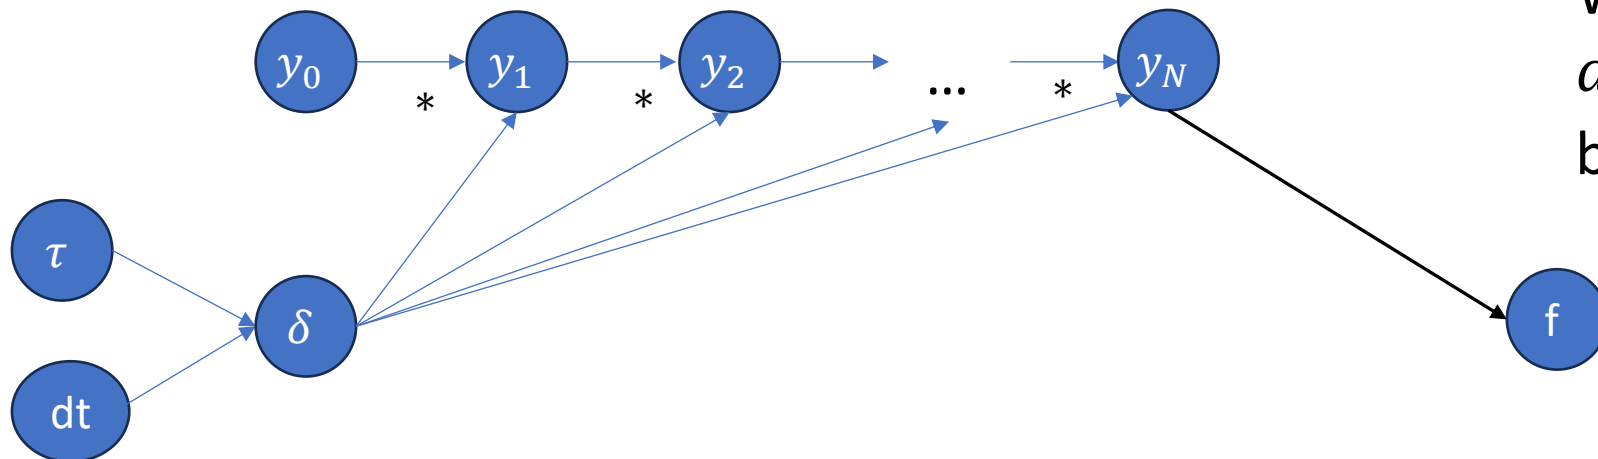
- For differential equations, the *adjoint equation* is itself another differential equation
- Backpropagation of gradients through our differential equation would implicitly solve *adjoint equations*
 - *We will show this...*

What if our function is a solution to an ODE?

- A simple example ODE, solved using Euler's method

$$\frac{dy}{dt} = -\frac{1}{\tau}y, \quad y_{n+1} = \left(1 - \frac{dt}{\tau}\right)y_n = \delta y_n$$

- Define some loss, f , which uses values of y
- It's just a graph



We can compute $d_\tau f$,
 $d_{dt} f$ and $d_{y_0} f$ using
back-propagation

What if our function is a solution to an ODE?

- A simple example ODE, with parameters y_0 and τ

$$\frac{dy}{dt} = -\frac{1}{\tau}y, \quad f = L(y(T))$$

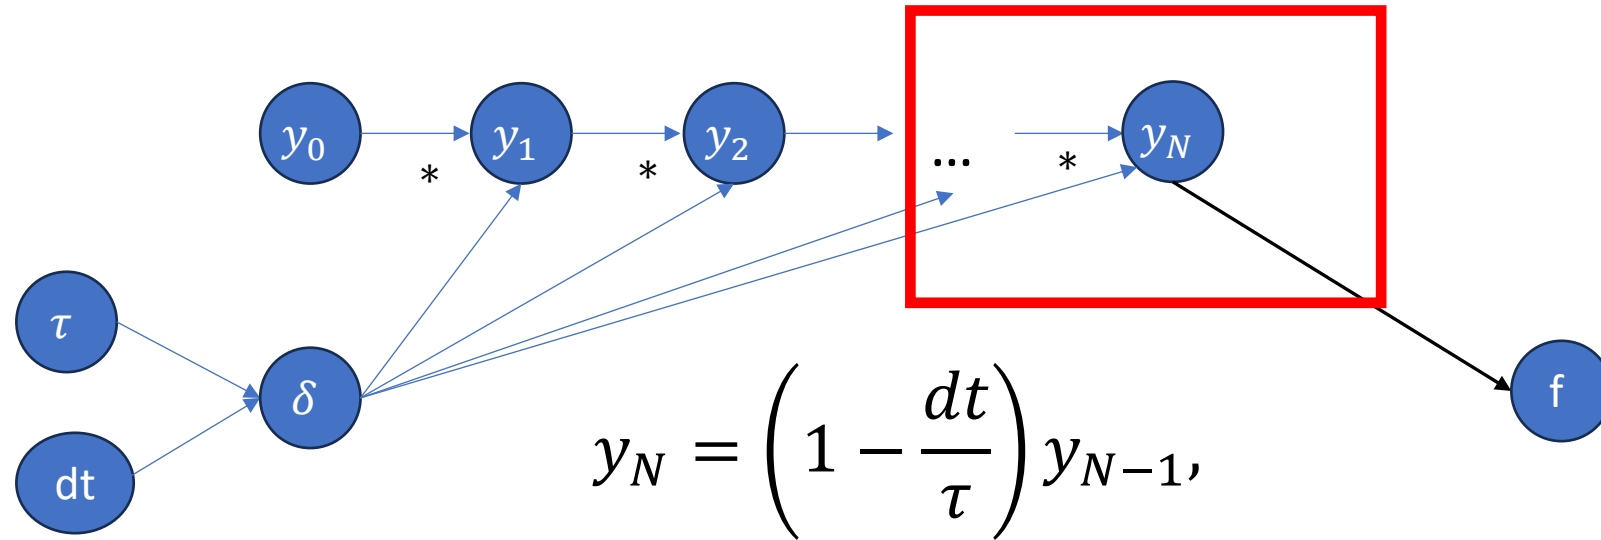
- What are the formal *adjoint equations*?

$$\text{Adjoint process: } \frac{d\alpha}{dt} = \frac{1}{\tau}\alpha, \quad \alpha(T) = \frac{dL}{dy(T)} \quad \text{Exponential growth}$$

$$\text{Parameter gradients: } \frac{d\beta}{dt} = -\frac{1}{\tau^2}\alpha y, \quad \beta(T) = 0$$

- Solved backwards in time, $\beta(0) = d_{\tau}f = \left(\frac{T}{\tau^2}\right) \left(\frac{dL}{dy(T)}\right) y_0 \exp\left[-\frac{T}{\tau}\right]$

What if our function is a solution to an ODE?

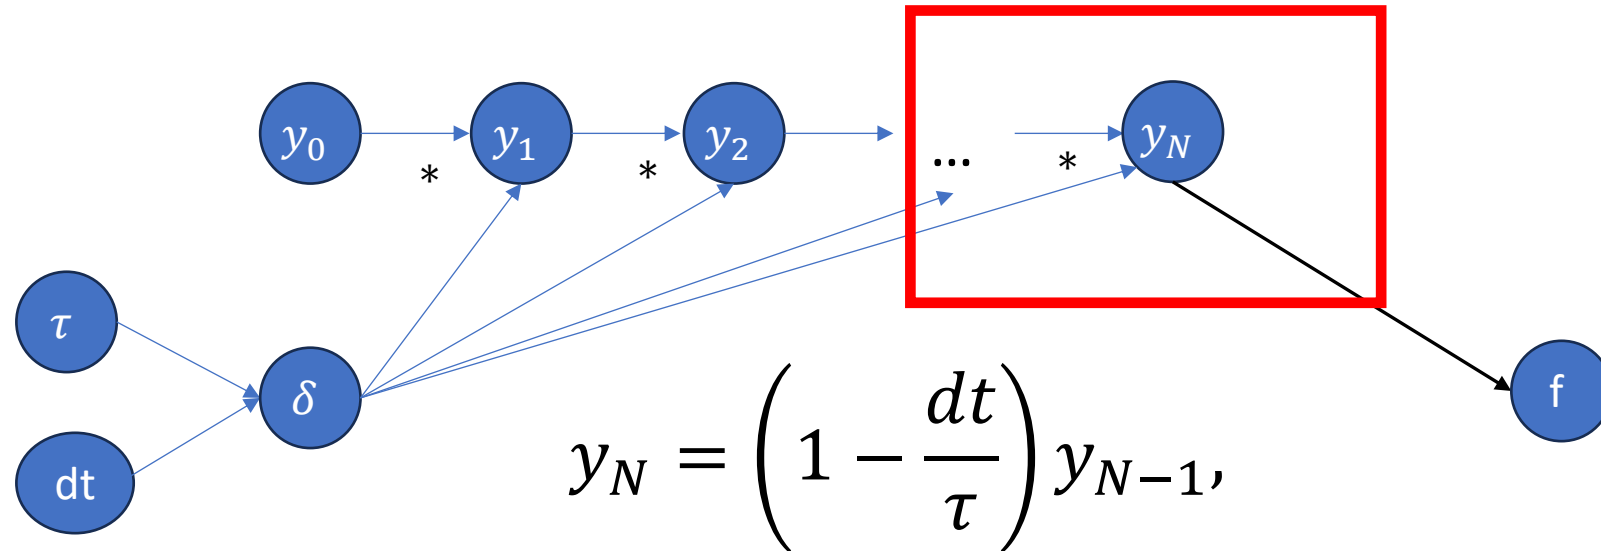


- Backpropagation a single step:

$$\frac{df}{d\tau} = \frac{dL}{dy_N} \cdot \frac{dy_N}{d\tau}$$

$$\frac{dy_N}{d\tau} = \frac{d}{d\tau} \left[\left(1 - \frac{dt}{\tau}\right) y_{N-1} \right]$$

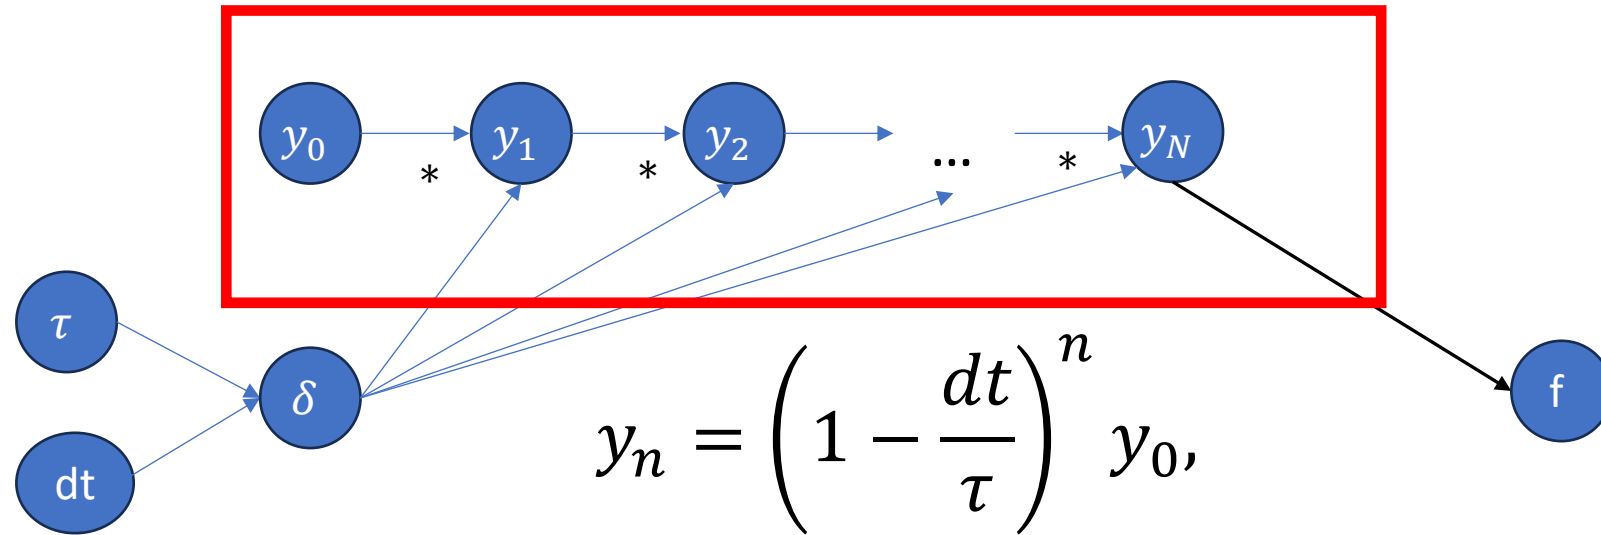
What if our function is a solution to an ODE?



- Backpropagation a single step:

$$\frac{df}{d\tau} = \frac{dL}{dy_N} \cdot \frac{dy_N}{d\tau} = \frac{dL}{dy_N} \cdot \left(\frac{dt}{\tau^2} y_{N-1} + \left(1 - \frac{dt}{\tau}\right) \frac{dy_{N-1}}{d\tau} \right)$$

What if our function is a solution to an ODE?



- Backpropagation all the way:

$$AD + ODE: \quad \frac{df}{d\tau} = \frac{T}{\tau^2} \frac{dL}{dy_N} y_0 \left(1 - \frac{T}{N\tau}\right)^{N-1}, \quad T = Ndt$$

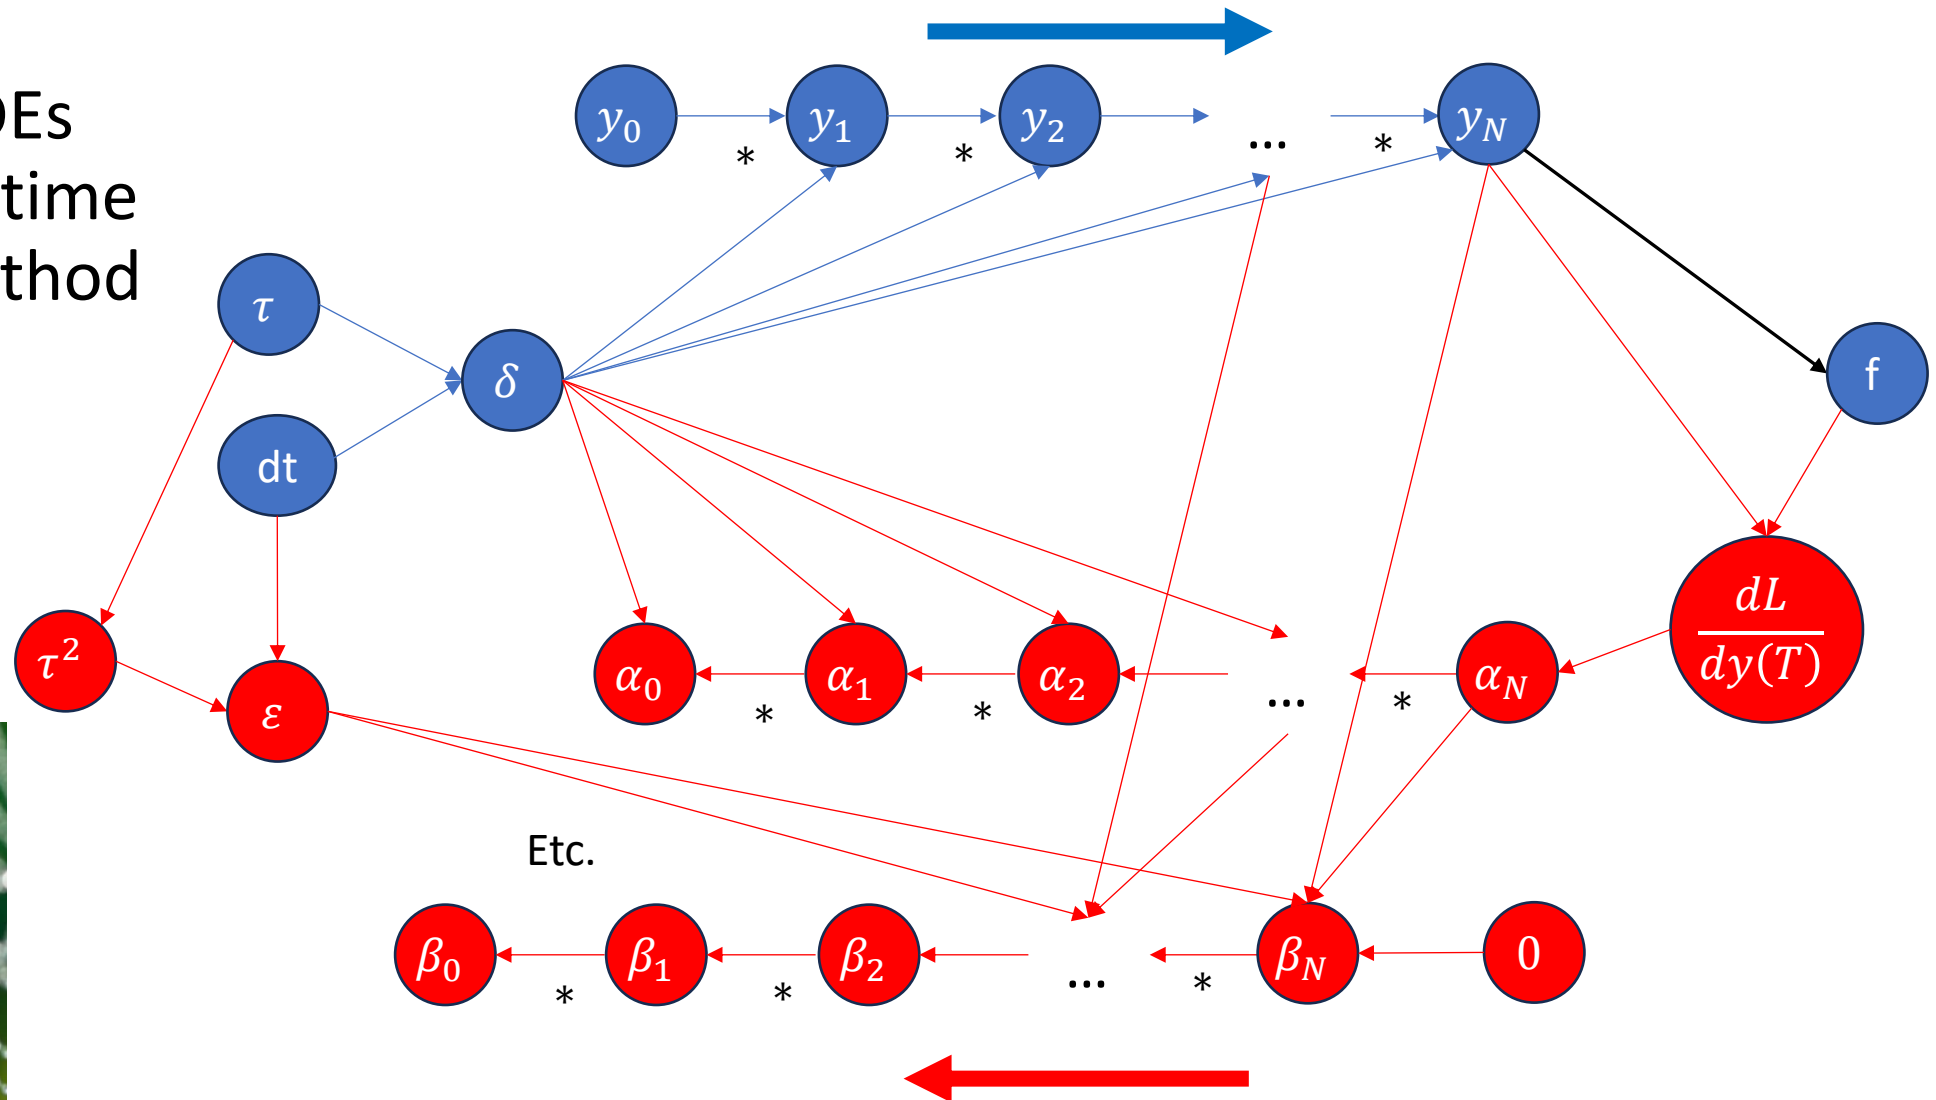
Analytic adjoint state method:

$$\beta(0) = \frac{T}{\tau^2} \frac{dL}{dy(T)} y_0 \exp \left[-\frac{T}{\tau} \right]$$

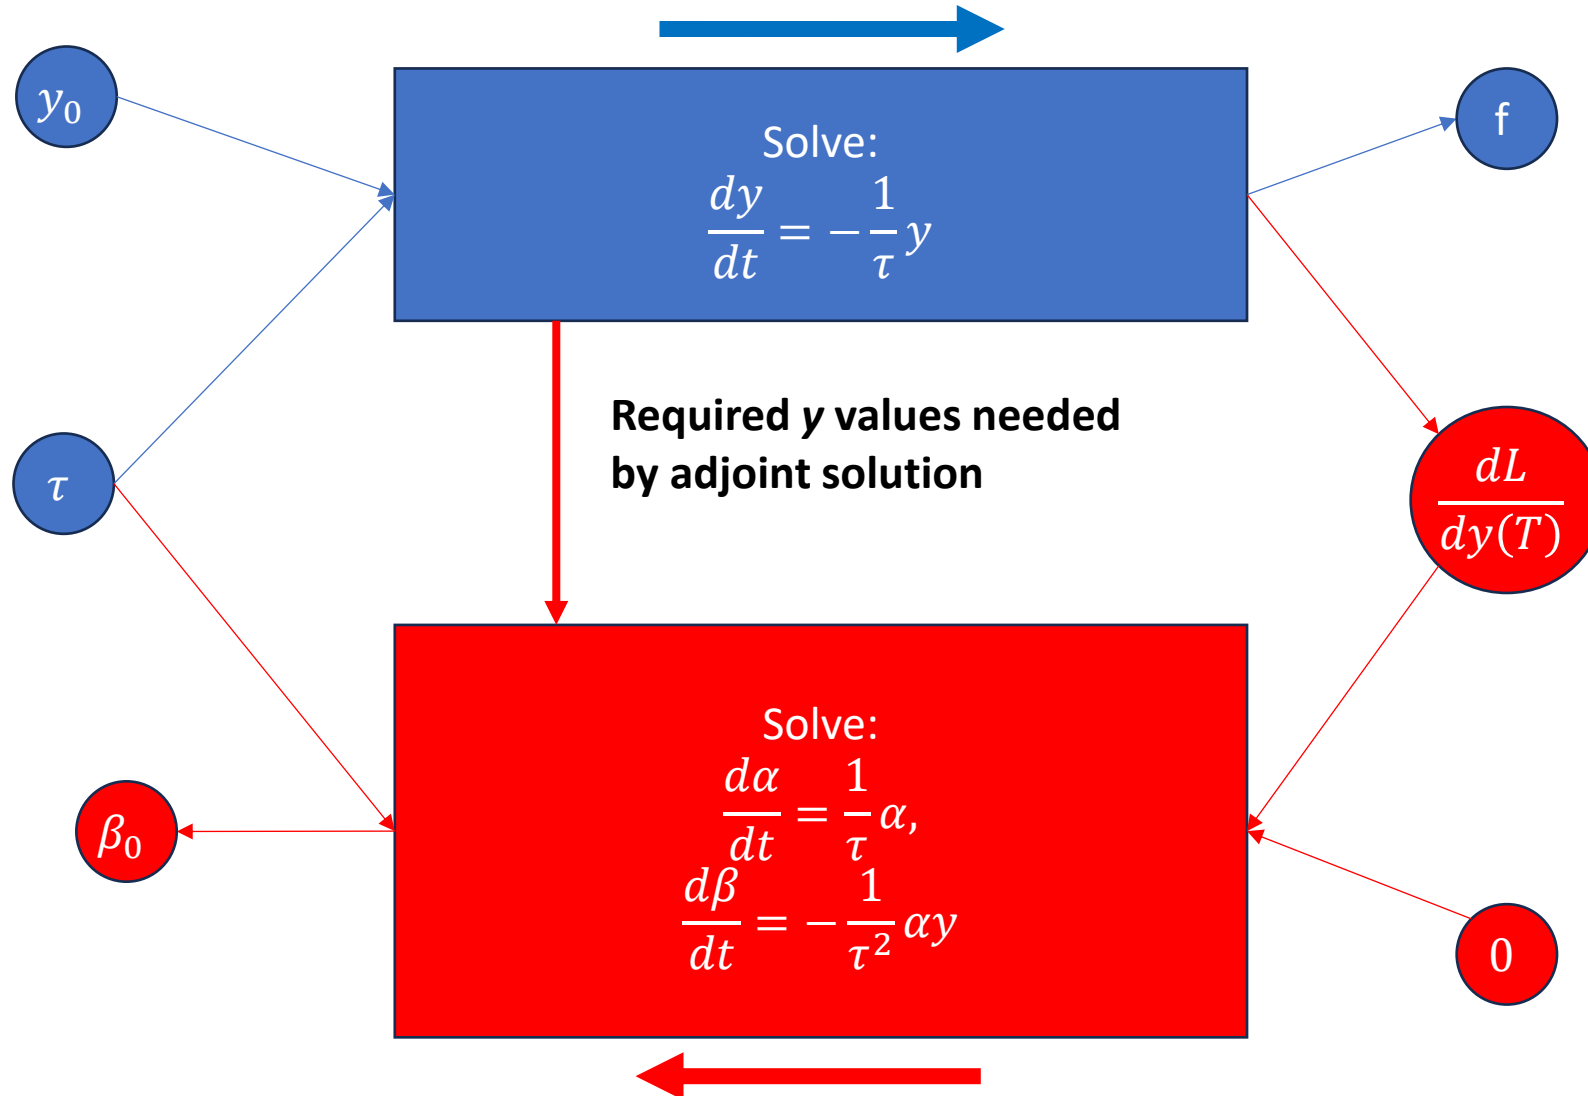
Backprop of ODE solution
numerically solves adjoint
state method
= “discretise-then-optimize”

Extending the computational graph

- Solving 3 ODEs with shared time step and method



An abstraction



Colab Exercise 1

- Computational graphs and AD
- Euler's method for simple ODEs
- The adjoint method for ODEs

*[https://github.com/
aidancrilly/MiniCourse-
DifferentiableSimulation](https://github.com/aidancrilly/MiniCourse-DifferentiableSimulation)*

Numerical modelling – PDEs

- Partial differential equations (PDEs) – dependent on only many independent variables
 - For the most part, these variables are space and time
- Total derivative vs. partial derivative

$$\frac{d}{dt} [f(x, t)] = \frac{\partial f}{\partial t} + \frac{\partial x}{\partial t} \cdot \frac{\partial f}{\partial x} = \frac{\partial f}{\partial t} + v_x \cdot \frac{\partial f}{\partial x}$$

- Example: Hydrodynamics

$$\begin{cases} \frac{\partial \rho}{\partial t} + \mathbf{u} \cdot \nabla \rho + \rho \nabla \cdot \mathbf{u} = 0 \\ \frac{\partial \mathbf{u}}{\partial t} + \mathbf{u} \cdot \nabla \mathbf{u} + \frac{\nabla p}{\rho} = \mathbf{g} \\ \frac{\partial e}{\partial t} + \mathbf{u} \cdot \nabla e + \frac{p}{\rho} \nabla \cdot \mathbf{u} = 0 \end{cases}$$

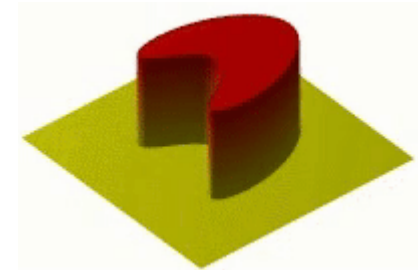


Numerical modelling – PDEs

- *Method of lines* for time dependent PDEs
 - Discretising all but the time dimension turns PDE → system of ODEs

- For example, the heat equation:

$$\frac{\partial T(x, t)}{\partial t} = k \frac{\partial^2 T(x, t)}{\partial x^2}$$



- Becomes (centred-space differencing):

$$\frac{dT(x_i, t)}{dt} = k \left(\frac{T(x_{i+1}, t) - 2T(x_i, t) - T(x_{i-1}, t)}{\Delta x^2} \right)$$

Numerical modelling – PDEs

- *Method of lines* for time dependent PDEs
 - Discretising all but the time dimension turns PDE → system of ODEs

- For example, the heat equation:

$$\frac{\partial T(x, t)}{\partial t} = k \frac{\partial^2 T(x, t)}{\partial x^2}$$

- Becomes:

$$\frac{d\mathbf{T}(t)}{dt} = k \underline{D} \mathbf{T}(t)$$

- Where \mathbf{T} is a vector of $T(x_i, t)$ at discrete points x_i

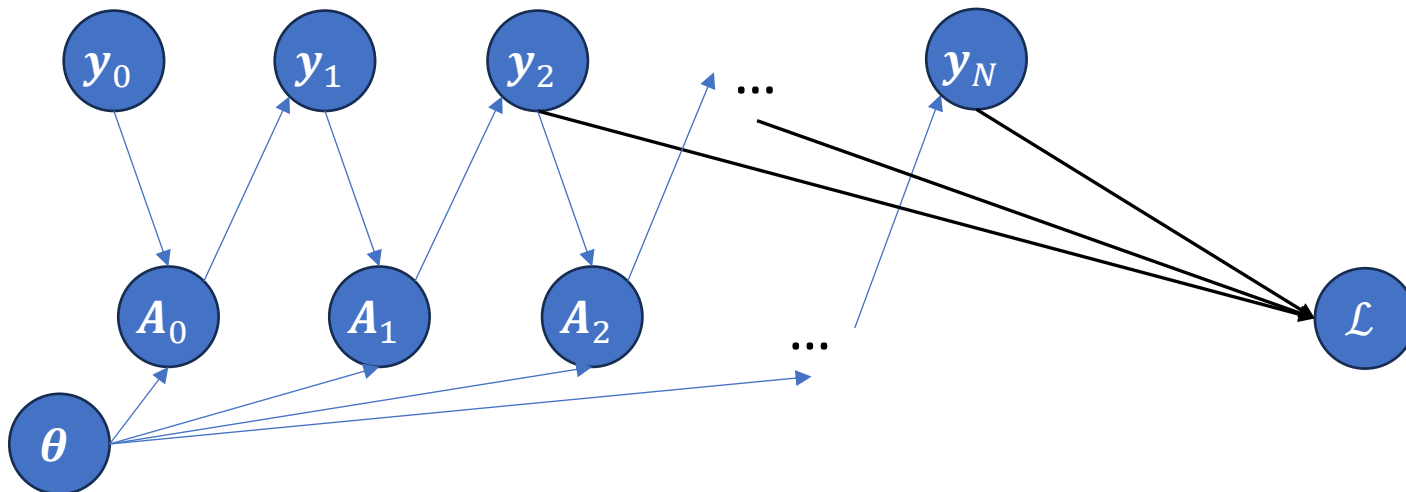
What if our function is a solution to a PDE?

- Spatial discretisation makes time-space PDEs into time ODEs
- Explicit time-stepping easiest:

$$\frac{d\mathbf{y}}{dt} = L(\mathbf{y}, \boldsymbol{\theta}) \xrightarrow{\text{discretise}} \mathbf{y}_{n+1} = (\mathbf{I} - \underline{L}(\mathbf{y}_n, \boldsymbol{\theta})dt)\mathbf{y}_n = \underline{A}_n\mathbf{y}_n$$

*Initial
conditions*

*Model
parameters*



We can modify our initial conditions or our model parameters to minimise the loss

Recipe for differentiable solver

1. Specify your ODE/PDE
2. Choose spatial (and temporal) differencing scheme
3. Define loss term for optimisation purposes
4. Compute adjoint system via 'discretise-then-optimize' or 'optimize-then-discretise' strategies

Gradient-Based Optimisation

- Key idea is gradient descent, move your parameters in the direction of local improvement
- Some other common concepts are:
 - *Momentum, \mathbf{m}* : The history of gradients is used to update parameters, not just current gradient value
 - *Adaptive learning rate, α* : The step size is adapted for each parameter to penalise parameters which cause oscillations in the gradient
- In general, the parameter updates look like:

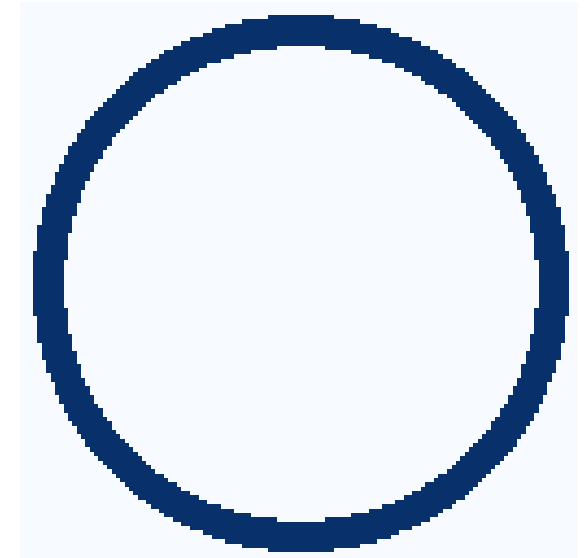
$$\boldsymbol{\theta}_{n+1} = \boldsymbol{\theta}_n + \alpha_n \mathbf{m}_n$$

Backpropagating through a fluid simulation

- Modify the initial conditions (velocities) to match some input image after a set simulation time

Steps:

1. Write a fluid simulator in a differentiable programming framework – in this case AutoGrad
2. Initialise some densities and x,y velocities on n_x by n_y grid
3. Find distance to target image (loss) after N simulation iterations
4. Use Jacobian of loss to update initial velocities ($2 \times n_x \times n_y$ values)



ODE differentiable programming libraries

- Diffrax – JAX-based library providing numerical differential equation solvers
- SciMLSensitivity.jl – Julia-based library, AD and adjoints for differential equations solvers (and more...)

ODE differentiable programming libraries

- **Difffrax – JAX-based library providing numerical differential equation solvers**
- SciMLSensitivity.jl – Julia-based library, AD and adjoints for differential equations solvers (and more...)

Difffrax example

- Difffrax has 4 key features (for our purposes):

1. Terms
2. Solvers
3. Adaptive stepping
4. Adjoint

```
def dydt(t,y,args):  
    return -y/args['tau']  
  
def difffrax_solve(dydt,t0,t1,Nt,rtol=1e-5,atol=1e-5):  
    """  
    Here we wrap the difffrax diffeqsolve function such that we can run with  
    different y0s and taus over the same time interval easily  
    """  
  
    # We convert our python function to a difffrax ODETerm  
    term = difffrax.ODETerm(dydt)  
    # We chose a solver (time-stepping) method from within difffrax library  
    # Heun's method (https://en.wikipedia.org/wiki/Heun%27s\_method)  
    solver = difffrax.Heun()  
  
    # At what time points you want to save the solution  
    saveat = difffrax.SaveAt(ts=jnp.linspace(t0,t1,Nt))  
    # Difffrax uses adaptive time stepping to gain accuracy within certain tolerances  
    stepsize_controller = difffrax.PIDController(rtol=rtol, atol=atol)  
  
    return lambda y0,tau : difffrax.diffeqsolve(term, solver,  
                                                y0=y0, args = {'tau' : tau},  
                                                t0=t0, t1=t1, dt0=(t1-t0)/Nt,  
                                                saveat=saveat, stepsize_controller=stepsize_controller)  
  
t0 = 0.0  
t1 = 1.0  
Nt = 100  
  
ODE_solve = difffrax_solve(dydt,t0,t1,Nt)  
  
# Solve for specific y0 and tau  
y0 = 1.0  
tau = 0.5  
sol = ODE_solve(y0,tau)
```

Difffrax example

- Difffrax has 4 key features (for our purposes):

1. Terms
2. Solvers
3. Adaptive stepping
4. Adjoint

```
def dydt(t,y,args):
    return -y/args['tau']

def difffrax_solve(dydt,t0,t1,Nt,rtol=1e-5,atol=1e-5):
    """
    Here we wrap the difffrax diffeqsolve function such that we can run with
    different y0s and taus over the same time interval easily
    """
    # We convert our python function to a difffrax ODETerm
    term = difffrax.ODETerm(dydt)
    # We chose a solver (time-stepping) method from within difffrax library
    # Heun's method (https://en.wikipedia.org/wiki/Heun%27s\_method)
    solver = difffrax.Heun()

    # At what time points you want to save the solution
    saveat = difffrax.SaveAt(ts=jnp.linspace(t0,t1,Nt))
    # Difffrax uses adaptive time stepping to gain accuracy within certain tolerances
    stepsize_controller = difffrax.PIDController(rtol=rtol, atol=atol)

    return lambda y0,tau : difffrax.diffeqsolve(term, solver,
                                                y0=y0, args = {'tau' : tau},
                                                t0=t0, t1=t1, dt0=(t1-t0)/Nt,
                                                saveat=saveat, stepsize_controller=stepsize_controller)

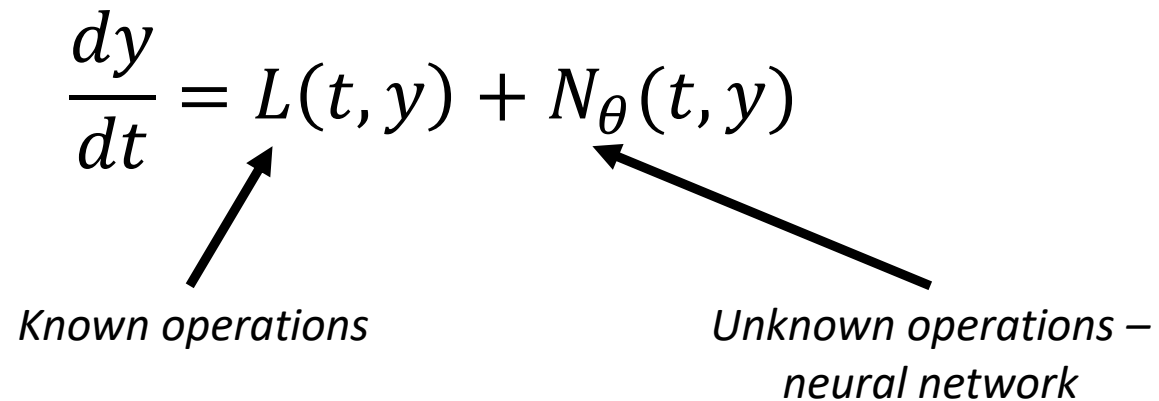
t0 = 0.0
t1 = 1.0
Nt = 100

ODE_solve = difffrax_solve(dydt,t0,t1,Nt)
def loss(inputs):
    y0 = inputs['y0']
    tau = inputs['tau']
    sol = ODE_solve(y0,tau)
    return sol.ys[-1]

inputs = {'y0' : y0, 'tau' : tau}
# Returns gradient of loss with respect to all inputs, i.e. dLdtau and dLdy0
jax.grad(loss)(inputs)
```

Neural Differential Equations

- What if what to describe a dynamical system but don't know the form of (some of) the terms?

$$\frac{dy}{dt} = L(t, y) + N_{\theta}(t, y)$$


Known operations

*Unknown operations –
neural network*

- Neural network parameters learnt using adjoint state method

Colab Exercise 2

- Diffrax library for solving ODEs
- Numerical solution to PDEs
- Differentiable PDE solvers
- Neural Differential Equations

*[https://github.com/
aidancrilly/MiniCourse-
DifferentiableSimulation](https://github.com/aidancrilly/MiniCourse-DifferentiableSimulation)*